


CMPT 606


Read Chapter 17

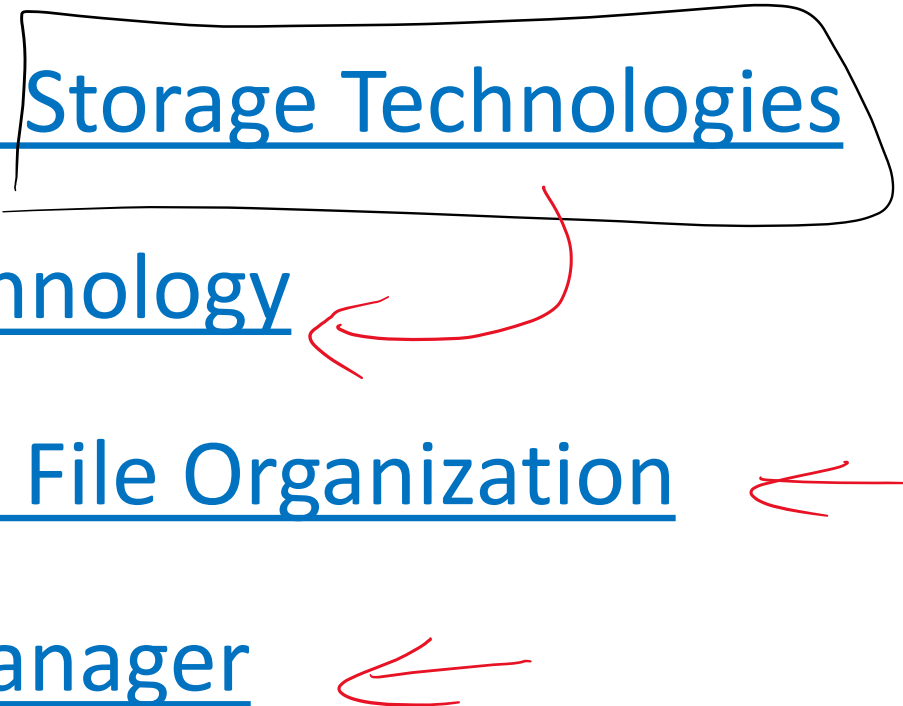
Database Storage

Dr. Abdelkarim Erradi

Department of Computer Science and Engineering

QU

Outline

- ① Database Storage Technologies
 - ② RAID Technology
 - ③ Database File Organization
 - ④ Buffer Manager
- 

Acknowledgment

Some slides are based on textbook slides &
CMU DB Course <https://15445.courses.cs.cmu.edu/fall2019/>

Course Outline – Database Inner Working

Storage

- **Disk Manager**
- **Buffer Pool Manager**
- Access Methods

Query Planning
and Execution


Concurrency
Control

Recovery

Disk-Oriented Architecture

- Disk-Oriented Architecture = primary storage location of the database is on non-volatile disk.
 - The DBMS's components manage the movement of data between Disk (non-volatile) and Memory (volatile).
- Storage Engine Design Goals
 - Allow the DBMS to manage **databases that exceed the available amount of memory**
 - Reading/writing to disk is expensive, so **I/O must be managed carefully** to avoid large waits and performance degradation

Database Storage Technologies



Why study the physical level of DBMS

- Someone has to write the DBMS software and its file manager!
- Some DB systems give the database administrator a range of physical storage options
 - Intelligent use of these options can make a very significant (and user-noticeable) difference in the way the system performs.
 - To "tune" the system properly, the DBA must understand what is happening at the physical level.
- Some of the techniques and algorithms can be used to solve other problems in other contexts

Key components of DBMS performance

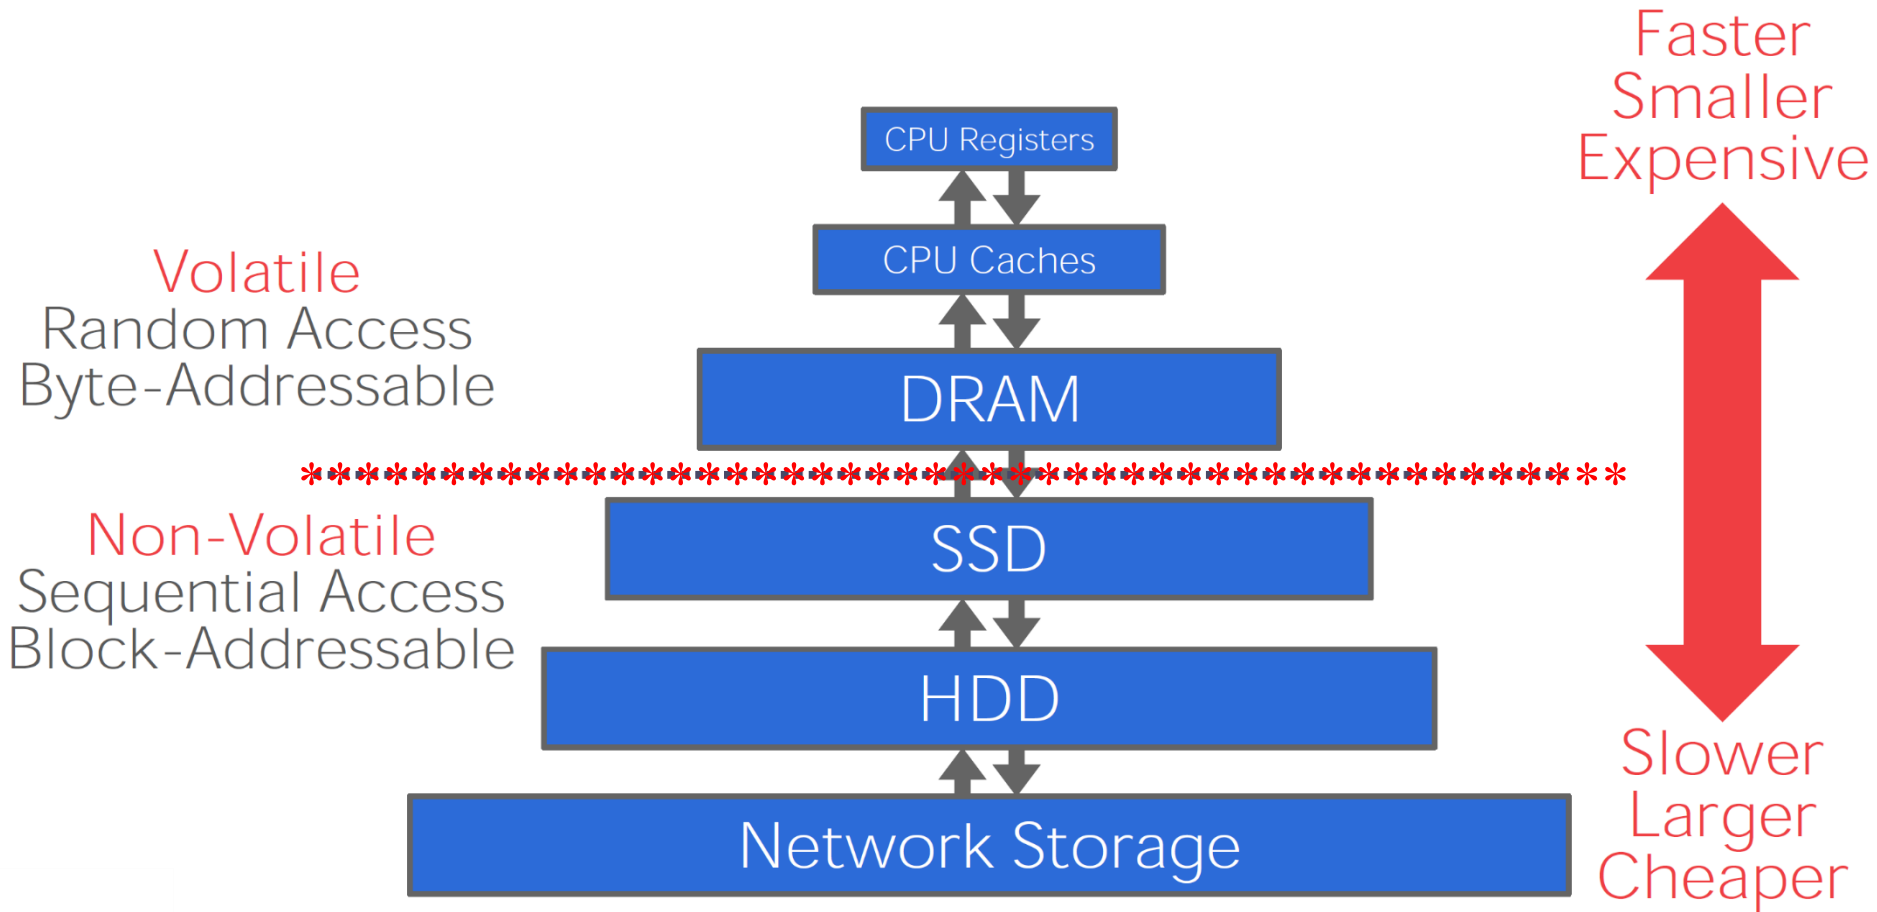
- The **performance of the DBMS Storage Engine** is often the key component of overall performance.
- There are two attributes that can be optimized:
 - 1. Response time** - defined as the time between the issuance of a command and the time that output for the command begins to be available.
 - e.g. if the command is a select statement, the time until the first row of the result appears

=> we want to minimize this

- 2. Throughput** - the number of operations that can be completed per unit time.

=> we want to maximize this

Storage Hierarchy



*** New trend: **Non-volatile memory** (e.g., Intel® Optane™)
<https://www.intel.com/content/www/us/en/architecture-and-technology/intel-optane-technology.html>

Storage Hierarchy - Primary storage

- **Primary storage:** Fastest media but volatile
 - Can hold subset of a database used by current transactions.
 - Volatile = data is lost when an application terminates (normally or due to a power failure or crash)
- **Cache**
 - Data and instructions in cache when needed by CPU.
 - On-board (L1) cache on same chip as CPU, L2 cache on separate chip.
 - Capacity couple of MBs, **access time few nanoseconds**
- **Main memory**
 - All active programs and data need to be in main memory.
 - Capacity couple of GBs, access time **10-100 nanoseconds**

Storage Hierarchy - Secondary & Tertiary Storage

- **Secondary storage:** non-volatile, moderately fast access time
 - Also called **online storage**
 - Stores the current version of entire database typically on a magnetic disk.
 - Access time in milliseconds
- **Tertiary storage:** non-volatile, slow access time
 - also called **off-line storage** – often used for archiving older versions of the database
 - Large capacity, access time seconds / minutes.
 - E.g. magnetic tape, optical storage

Large speed gap between Memory and Disk

- The large speed gap between Memory and Disk remains the key issue in DBMS performance.
- Time to access information in disk is the **major determining factor in system performance**.
- The *number of disk I/Os* (block accesses) is often used to measure the cost of a database operation.

Relative Daps in Access Time

| | | |
|------------------|-----------------|---------------|
| 0.5 ns | L1 Cache Ref | ← 0.5 sec |
| 7 ns | L2 Cache Ref | ← 7 sec |
| 100 ns | DRAM | ← 100 sec |
| 150,000 ns | SSD | ← 1.7 days |
| 10,000,000 ns | HDD | ← 16.5 weeks |
| ~30,000,000 ns | Network Storage | ← 11.4 months |
| 1,000,000,000 ns | Tape Archives | ← 31.7 years |

Source: <https://gist.github.com/hellerbarde/2843375>

- Each level is thousands of times faster than the level below it.
- **Dominance of I/O cost:** A modern CPU can execute millions of instructions while reading a block.

Hard Disks

- Secondary storage device of choice.
- Data is stored and retrieved in units called *disk blocks* or *pages* (typically 4 or 16 kilobytes)
- Unlike RAM, time to retrieve a disk page varies depending upon location on disk.
 - Reading several pages in sequence from a disk takes much less time than reading several random pages
- Therefore, **relative placement of pages** on disk has major impact on DBMS performance!

What's Inside A Disk Drive?





Image courtesy of Seagate Technology

Inside the HD <http://www.youtube.com/watch?v=kdmLv11n82U>



Disk Physical Structure

- Disks consist of **platters**, each with two surfaces
- Each surface consists of concentric rings called **tracks**
- Each track consists of **sectors** separated by gaps
 - Track capacities vary typically from 4 to 50 Kbytes or more
- All tracks under heads at the same time make a **cylinder** (imaginary!).
- Only one head reads/writes at any one time.




Disk Operation (Single-Platter View)

- The disk surface spins at a fixed rotational rate
- The read/write head is attached to the end of the arm and flies over the disk surface
- By **moving radially**, the arm can position the read/write head over any track




Disk Access – Service Time Components




Seek for **RED**

↑
Seek
moving arms to
position disk head
on track



Rotational latency

↑
Rotational latency
waiting for block to
rotate under head



After **RED** read

↑
Data transfer
moving data to/
from disk surface

Typically about 1% of the time is actually spent on data transfer, the rest is access time.

Performance Measures of Disks

- **Access time** – the time it takes from when a read or write request is issued to when data transfer begins. Consists of:
 - **Seek time** – time it takes to reposition the arm over the correct track.
 - 4 to 10 milliseconds on typical disks
 - **Rotational latency** – time it takes for the sector to be accessed to appear under the head.
 - 4 to 11 milliseconds on typical disks (5400 to 15000 r.p.m.)
 - **Data-transfer rate** – the rate at which data can be retrieved from or stored to the disk.
 - 25 to 100 MB per second rate, lower for inner tracks
- Access time dominated by **seek time** and **rotational latency**.
 - Disk is about 40,000 times slower than RAM

Disk speeds are dominated by access time

- For this reason, information on disk is always organized in Blocks – blocks are **basic units of transfer and storage**
 - relatively large chunks (4 or 16 kilobytes) of contiguous information that is read/written as a unit.
 - it always **reads or writes the whole block** containing a desired piece of information.
 - A system never reads or writes a single disk byte.
 - The block size B is fixed for each system.
 - Typical block sizes range from 4 to 16 kilobytes
- Mapping between logical blocks and actual (physical) sectors
 - Maintained by hardware/firmware device called disk controller.
 - Converts requests for logical blocks into (surface,track,sector) triples.

Optimization of Disk Block Access

A major goal of the design of DBMS file systems is to **minimize the time spent waiting for disk accesses**. 3 ways this is done:

- 1) **Store related information on the same or nearby blocks:**
read and write of data on ***contiguous disk blocks*** and eliminates seek time and rotational delay time for all but the first block transfer
 - Files may get fragmented over time (if data is inserted to/deleted from the file) => reorganize the database files to speed up access
- 2) **Keeping copies of recently-used information in buffers in memory**, so that if the same information is needed again it can be accessed without having to go to the disk again
- 3) **Parallelism** - spreading information across multiple disks, so that several disks can be going through the physical operations needed to access information at the same time

Example: reading two disk blocks

- Assume
 - average seek time = 10 ms
 - average rotational latency = 3 ms
 - transfer time for 1 block = 0.01875 ms
- **Adjacent** block on same track
 - access time = $10 + 3 + 2 \times (0.01875)$ ms = 13.0375 ms
- **Random** block
 - access time = $2 \times (10 + 3 + 0.01875)$ ms = 26.0375 ms

RAID Technology

Parallelizing Disk Access using RAID Technology

- **RAID: Redundant Arrays of Independent Disks**
 - an array of independent disks **acting as a single higher-performance logical disk**, providing:
 - **high capacity** and **high speed** by using multiple disks in parallel,
 - **high reliability** by storing data redundantly, so that data can be recovered even if a disk fails
- The main goal of RAID is to **reduce the large speed gap between disks and the memory**

RAID goals




RAID systems seek :

- **to improve throughput** by a technique known as **striping**, in which a single file is spread over multiple disks.
 - **Parallelize** large accesses to reduce response time: multiple accesses to different parts of the same file can often be performed in parallel (assuming that the parts being accessed are on different disks).
- **to improve reliability** by **replication** of data, so that if a disk fails, the data it contained is available somewhere else.
 - => improve throughput for reads -if there are multiple copies of an item, then any copy can be read.
 - but creates an issue on write though - since all copies must be updated.

RAID 0 - a.k.a. Striping

- Requires two or more disks
- No lost drive space due to striping
- Fastest read and write performance.
- Raid level 0 has no redundant data and hence has the best write performance at the risk of data loss
 - Offers no data protection.
- The more disks, the more risk.

Used in high-performance applications where data loss is not critical



Data striping

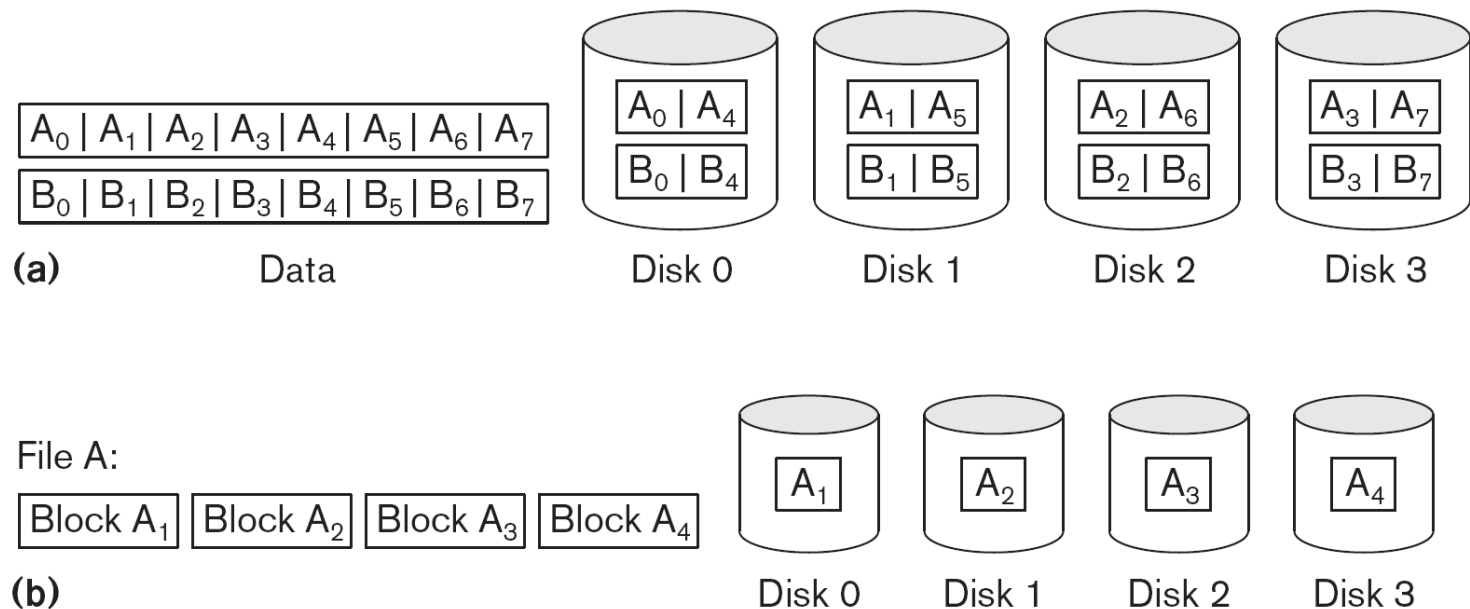
- RAID uses a concept called **data striping** = distribute blocks over n disks in a round robin fashion.
 - Make disks appear as a single large, fast disk.
 - Requests for different blocks can run in parallel if the blocks reside on different disks

Figure 17.13

Striping of data across multiple disks.

(a) Bit-level striping across four disks.


(b) Block-level striping across four disks.



RAID 1 - a.k.a. Mirroring


- Raid level 1 uses mirrored disks
- Write speed of one disk
- Read speed of two disks
- Capacity is equal to the size of one

Popular for applications such as storing log files in a database system.




RAID 5 - Striping with Distributed Parity

- Considered best compromise between speed and storage efficiency:
 - Good performance (as blocks are striped) but slower writes due to parity
 - Good redundancy (distributed parity)
- Requires 3 or more drives
- Stripe across all drives with **parity**
- Can loose 1 drive and still function
- Capacity is **$n-1$** where **n** is number of drives in array



RAID 6 - RAID 5 on Steroids

- 4 or more disk
- Is a stripe with two parity drives
- Can loose two drives and still function
- Capacity is $n-2$ where n is number of drives in array
- Protect against up to two disk failures by using just two redundant disks



Comparison of Single RAID Levels

| | RAID 0 | RAID 1 | RAID 5 | RAID 6 |
|-----------------------|--------------------------------|----------------|--------------------------------------|--------------------------------------|
| Diagram | | | | |
| Description | Striping | Mirroring | Striping with Parity | Striping with Dual Parity |
| Minimum Disks | 2 | 2 | 3 | 4 |
| Array Capacity | No. of Drives x Drive Capacity | Drive Capacity | (No. of Drives - 1) x Drive Capacity | (No. of Drives - 2) x Drive Capacity |

Comparison of Single RAID Levels

| | RAID 0 | RAID 1 | RAID 5 | RAID 6 |
|--|--------|--|---|--|
| Storage Efficiency | 100% | 50% | (Num of drives – 1) / Num of drives | (Num of drives – 2) / Num of drives |
| Fault Tolerance | None | 1 Drive failure | 1 Drive failure | 2 Drive failures |
| High Availability | None | Good | Good | Very Good |
| Degradation during <u>rebuild</u> | NA | <ul style="list-style-type: none"> • Slight degradation • Rebuilds very fast | <ul style="list-style-type: none"> • High degradation • Slow Rebuild (due to write penalty of parity) | <ul style="list-style-type: none"> • Very High degradation • Very Slow Rebuild (due to write penalty of dual parity) |


Understanding the Parity

- RAID 5 and RAID 6 store parity information against data for rebuild
- Parity can be calculated using a simple XOR
- eg– “ABCDEFGHijkl” on a 4 disk RAID 5 array

| Disk 1 | Disk 2 | Disk 3 | Disk 4 |
|--------------|--------------|--------------|----------------|
| A (01000001) | B (01000010) | C (01000011) | {P – 01000000} |
| Parity {P} | D | E | F |
| G | Parity {P} | H | I |
| J | K | Parity {P} | L |

- If Disk 2 fails then the data “B” can be recalculated as
(01000001 XOR 01000011 XOR 01000000) => 01000010 => B
- More info @ <http://www.youtube.com/watch?v=LTq4pGZtzh0>

RAID 10 a.k.a. 1+0



Diagram

Description

Mirroring then Striping

Minimum Disks

Even number > 4

Maximum Disks

Controller Dependant

Array Capacity

$(\text{Size of Drive}) * (\text{Number of Drives}) / 2$

Storage Efficiency

50%

Fault Tolerance

Multiple drive failure as long as 2 drives from same RAID 1 set do not fail


High Availability

Excellent

Choice of RAID Level

- Factors in choosing RAID level
 - Monetary cost
 - Performance: Number of I/O operations per second, and bandwidth during normal operation
 - Performance during failure
 - Performance during rebuild of failed disk
- RAID 0 is used only when data safety is not important
 - E.g. data can be recovered quickly from other sources
- Level 6 is rarely used since levels 1 and 5 offer adequate safety for most applications


Networked Storage



Source: <http://www.youtube.com/watch?v=2T99tW1KEMc>

Storage Area Network (SAN)


- Online storage peripherals are configured as nodes on a **high-speed network** and can be attached and detached from servers in a very flexible manner
- Servers see SAN as a virtual drives
- **Dedicated access** - each part of the SAN is dedicated to each server
- **Block based storage**



- Storage traffic segregating from the rest of the LAN. It typically uses Fiber Channel connectivity

Network Attached Storage (NAS)

- File Server optimized to serve files over the main LAN (OS dedicated to file system)
- **File based storage**
- Servers see NAS as a Network Share (need to map it to a drive)
- Suitable for sharing files



Database File Organization


Database Storage

Database Storage addresses 2 problems:

- **Problem #1:**How the DBMS represents the database in files on disk.
 - File Storage
 - Page Layout
 - Record Layout
- **Problem #2:**How the DBMS manages its memory and efficiently move data back-and-forth from disk.

File Organization

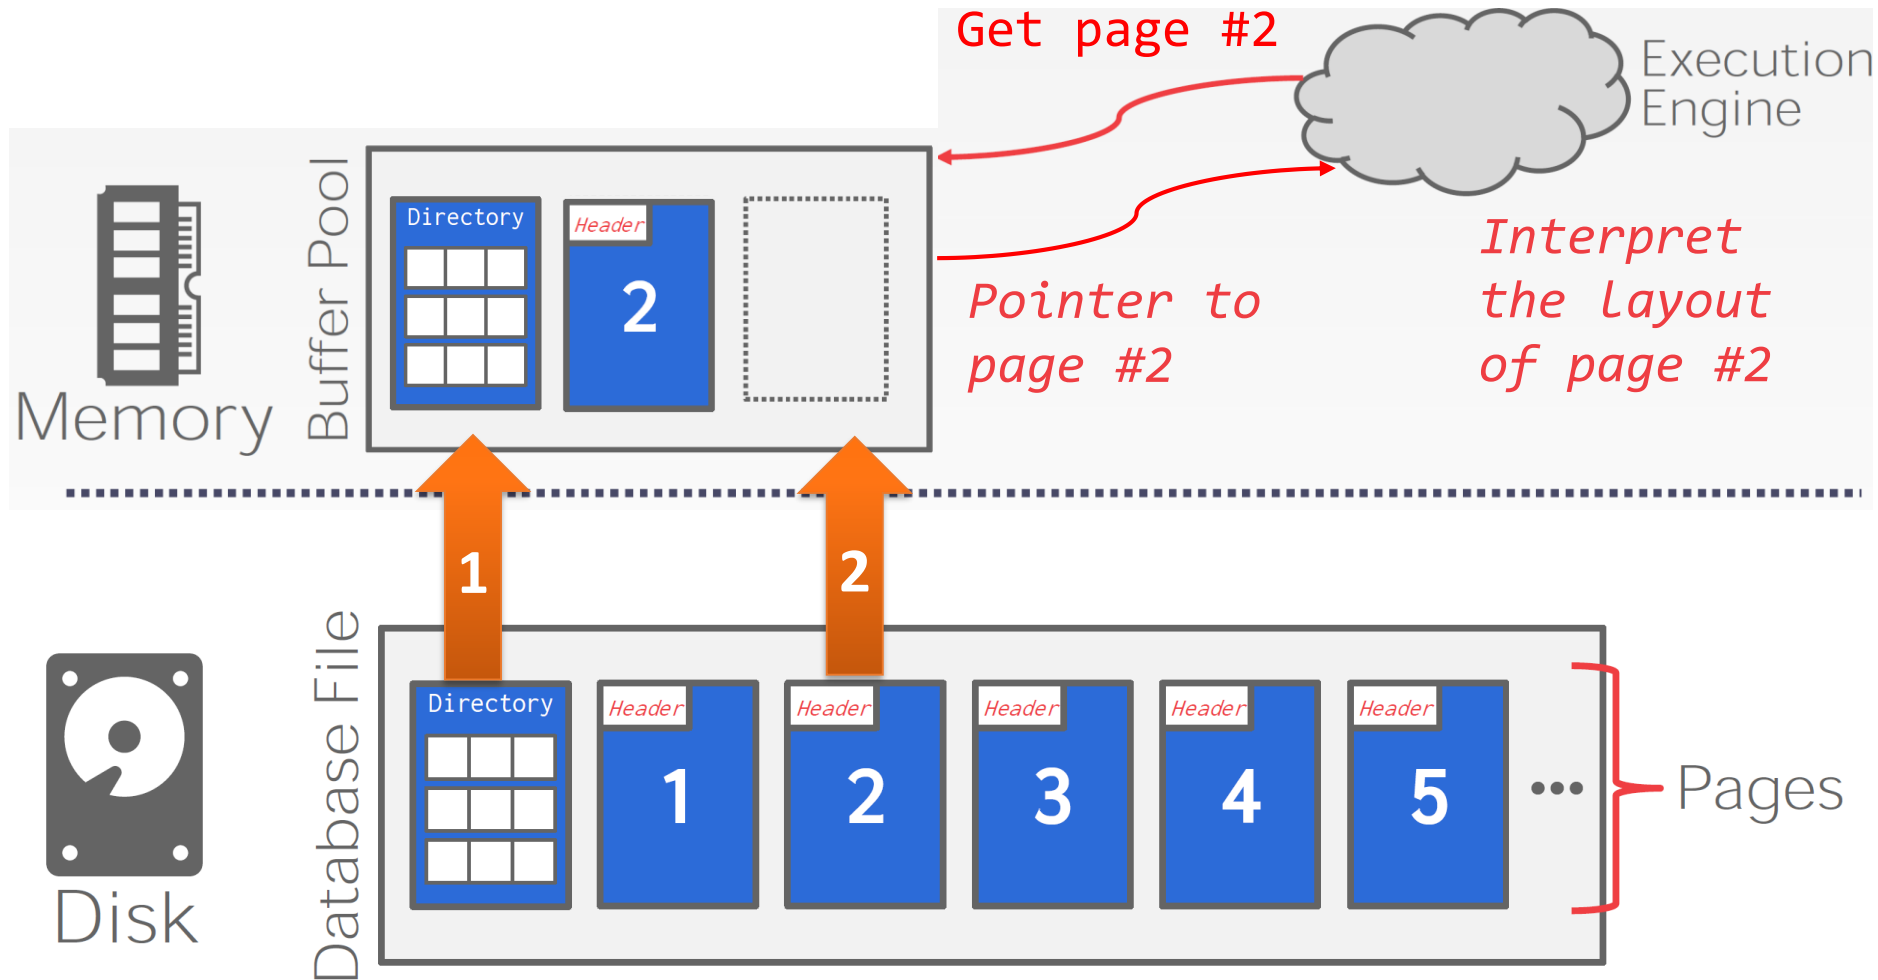
- The DBMS stores a database as one or more *files* on disk.
 - The OS doesn't know anything about the contents of these files.
- Each file has a collection of *pages* (aka blocks)
- Each *page* is a sequence of *records*.
- A record is a sequence of *fields*.



Storage Manager

- The **storage manager** is responsible for maintaining a database's files.
 - It organizes the files as a collection of pages.
 - Tracks data read/written to pages.
 - Tracks the available space.

How the Storage Engine Works?









Why NOT use the OS?

- In DBMS the OS is NOT used for moving data pages in and out of memory.
- DBMS (almost) always wants to control things itself and can do a better job at it.
 - Flushing dirty pages to disk in the correct order.
 - Specialized prefetching.
 - Buffer replacement policy.
 - Thread/process scheduling.

Database Pages

- A **page** is a **fixed-size** block of data (4 to 16KB).
 - It can contain records, meta-data, indexes, log records...
- Each page is given a unique identifier.
- The DBMS uses an indirection layer (i.e., a **dictionary**) to map page ids to physical locations.


| | |
|-------|--|
| 4KB |  |
| |  |
| |  |
| <hr/> | |
| 8KB |  |
| |  |
| <hr/> | |
| 16KB |  |

Heap File

- DBMSs manage pages in files on disk in different ways:
 - Heap File Organization
 - Sequential / Sorted File Organization
 - Hashing File Organization
- A **heap** file is an unordered collection of pages where records are stored in random order. Supported operations include:
 - Create / Get / Write / Delete Page
 - Iterate over all pages
- Need meta-data to keep track of what pages exist and which ones have free space. Typically using a **Page Directory**

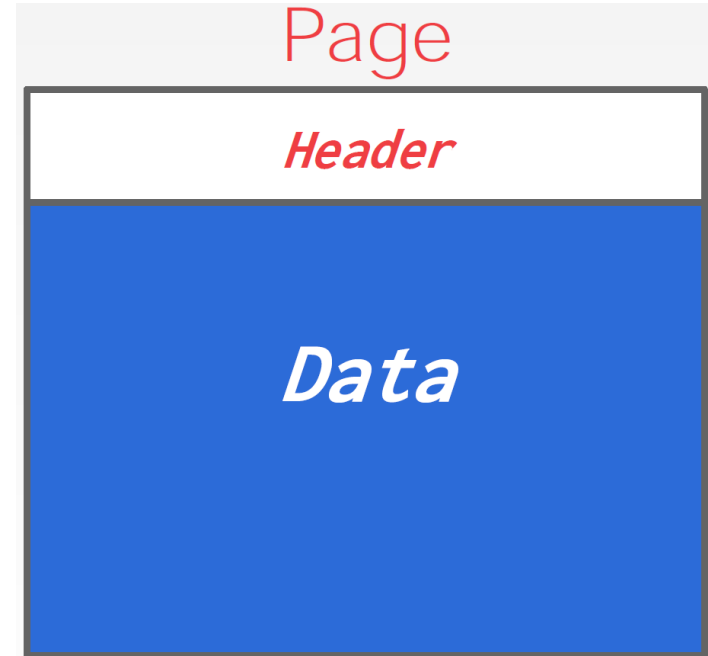
Heap File – Page Directory

- The DBMS maintains special pages that **tracks the location of data pages** in the database files.
- The directory also records the **number of free slots** per page.
- The DBMS has to make sure that the directory pages are in sync with the data pages.




Page Layout

- Every page contains a **header** of meta-data about the page's contents:
 - Page Size
 - Checksum
 - DBMS Version
 - Compression Information
- The page data is organized using 3 approaches:
 - Row-oriented
 - Column-oriented
 - Log-structured




Row-oriented Storage

- The most common layout scheme is called **slotted pages**.
- The slot array maps "slots" to the records' **starting position** offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.



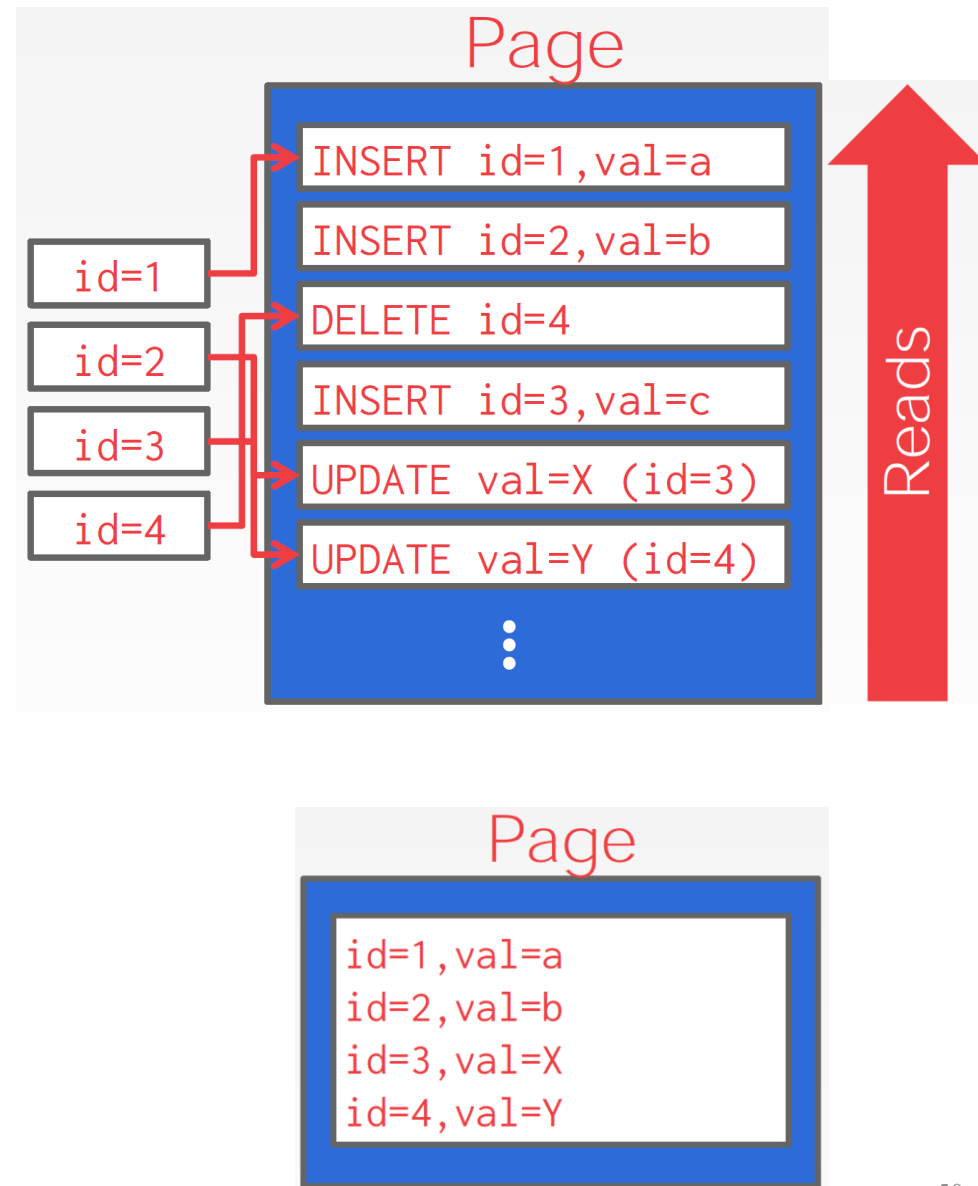
Log-structured Organization

- Instead of storing records in pages, the DBMS only stores **log records**.
- The system appends log records to the file of how the database was modified:
 - Inserts store the entire tuple.
 - Deletes mark the tuple as deleted.
 - Updates contain the **delta** of just the attributes that were modified.



Log-structured Organization

- To read a record, the DBMS scans the log **backwards** and "recreates" the record
- Indexes can be used to allow it to jump to locations in the log.
- Periodically **compact** the log.

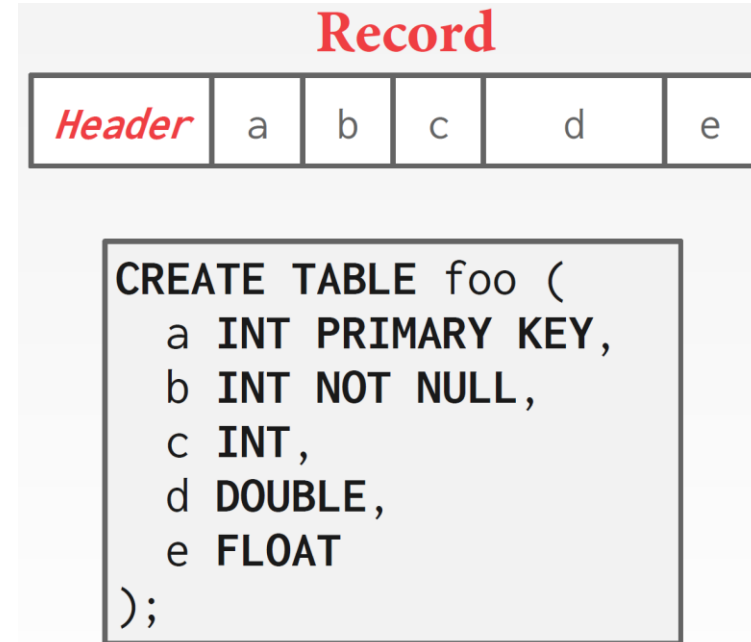


Record Layout

- A record is essentially a sequence of bytes.
- It's the job of the DBMS to interpret those bytes into attribute **types** and **values**.
- DBMS's **catalog** contain the schema information about tables that the system uses to figure out the record's layout.
- The DBMS assigns each record a **unique record identifier** to keep track of individual records

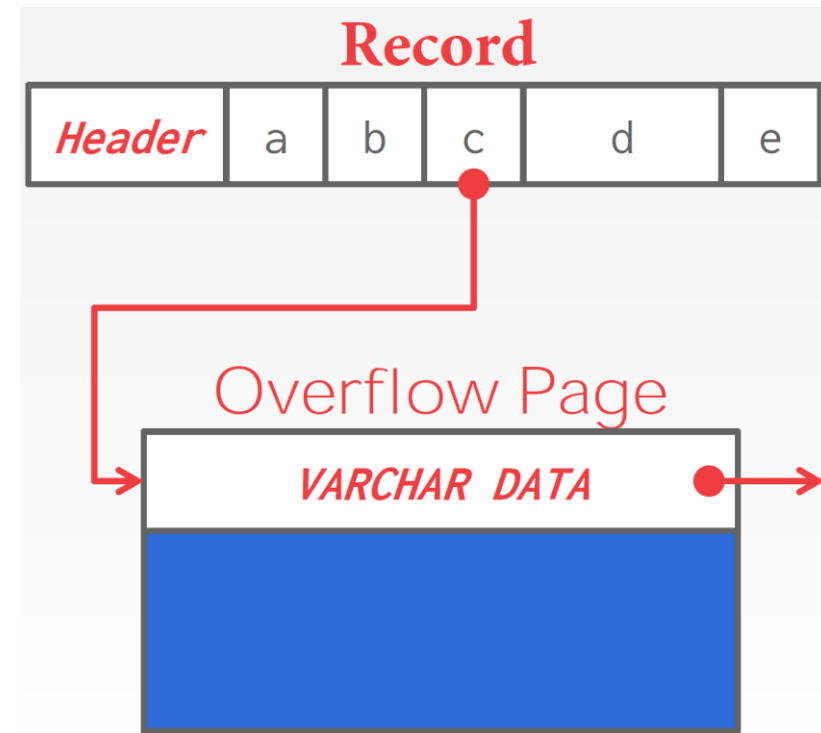
Most common:

Record id = page_id + slot



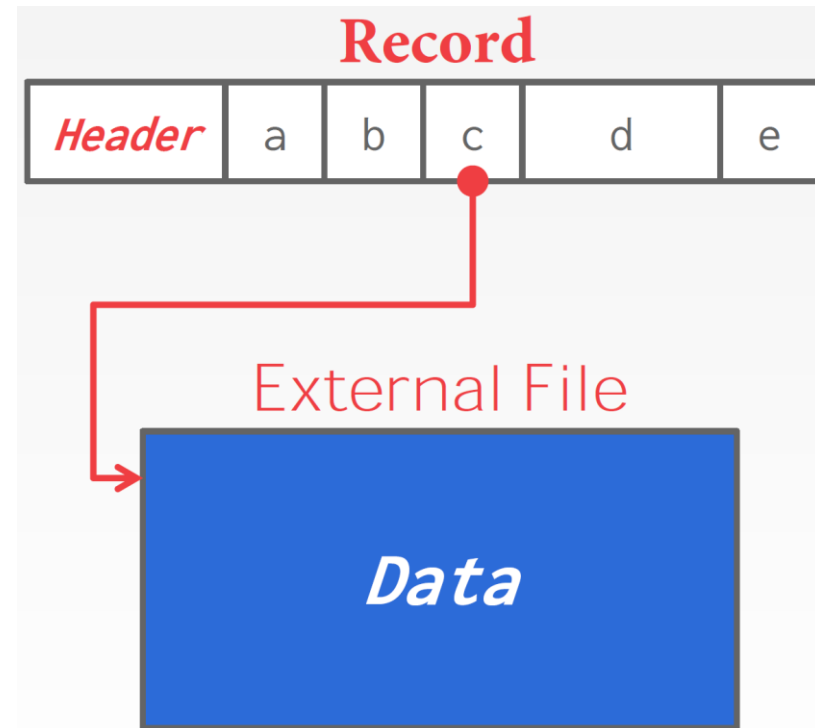
Storing Large Values

- Most DBMSs don't allow a record to exceed the size of a single page.
- To store values that are larger than a page, the DBMS uses separate **overflow storage pages**.



External Value Storage

- Some systems allow you to store a really large value in an external file. Treated as a **BLOB** type.
 - Oracle: **BFILE** data type
 - Microsoft: **FILESTREAM** data type
- The DBMS **cannot** manipulate the contents of an external file
 - No durability protections
 - No transaction protections



OLTP vs. OLAP

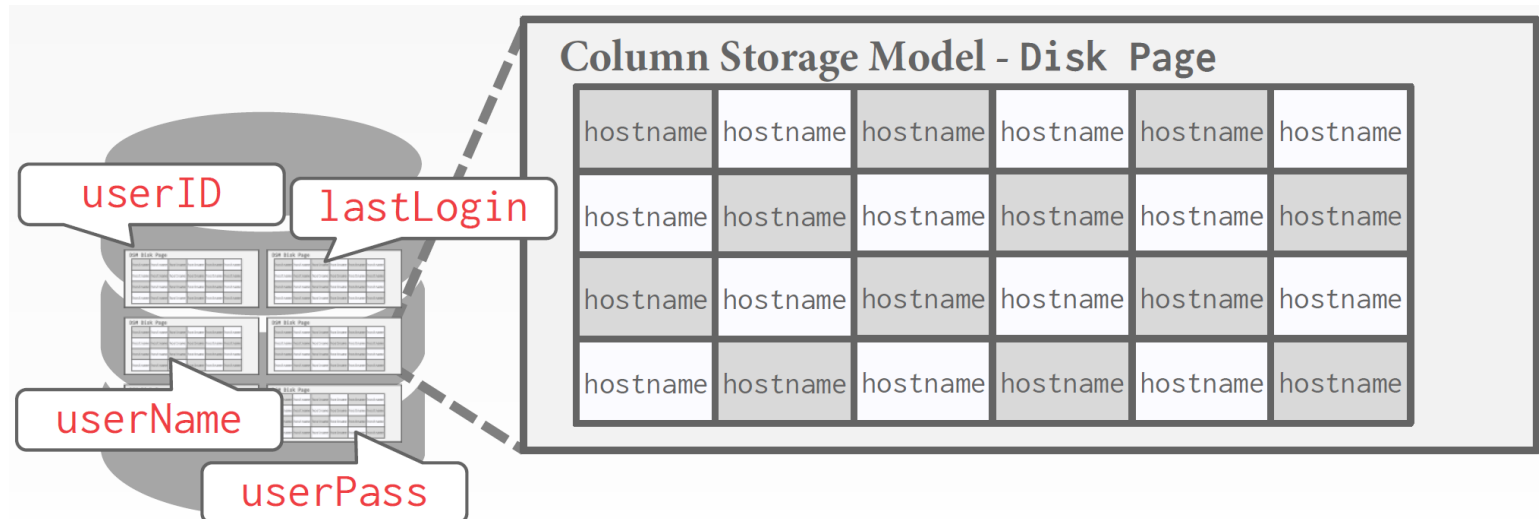
- On-line Transaction Processing:
 - Simple queries that read/update a small amount of data that is related to a single entity in the database.
- On-line Analytical Processing:
 - Complex queries that read large portions of the database spanning multiple entities.
- The DBMS can store records in different ways that are better for either OLTP or OLAP workloads:
 - OLTP = Row Store
 - OLAP = Column Store

Row storage model

- Row storage model (aka "n-ary storage model")
- The DBMS stores all attributes for a single record contiguously in a page.
- Ideal for OLTP workloads where queries tend to operate only on an individual entity and insert heavy workloads.
- Advantages
 - Fast inserts, updates, and deletes.
 - Good for queries that need the entire tuple.
- Disadvantages
 - Not good for scanning large portions of the table and/or a subset of the attributes.

Column Storage Model

- The DBMS stores the values of a single attribute for all records contiguously in a page.
- Ideal for OLAP workloads where read-only queries perform large scans over a subset of the table's attributes.



Column Storage Model

- **Advantages**

- Reduces the amount wasted I/O because the DBMS only reads the data that it needs.
- Better query processing and data compression

- **Disadvantages**

- Slow for point queries, inserts, updates, and deletes because of tuple splitting/stitching.