

CMPT 606

Read Chapter 17

Database Storage



Dr. Abdelkarim Erradi

Department of Computer Science and Engineering

QU

Outline

- ① Database Storage Technologies
- ② RAID Technology
- ③ Database File Organization
- ④ Buffer Manager

Acknowledgment

Some slides are based on textbook slides &
CMU DB Course <https://15445.courses.cs.cmu.edu/fall2019/>

Course Outline – Database Inner Working

Storage

- **Disk Manager**
- **Buffer Pool Manager**
- Access Methods

Query Planning
and Execution


Concurrency
Control

Recovery

Disk-Oriented Architecture

- Disk-Oriented Architecture = primary storage location of the database is on non-volatile disk.
 - The DBMS's components manage the movement of data between Disk (non-volatile) and Memory (volatile).
- Storage Engine Design Goals
 - Allow the DBMS to manage **databases that exceed the available amount of memory**
 - Reading/writing to disk is expensive, so **I/O must be managed carefully** to avoid large waits and performance degradation

Database Storage Technologies



Why study the physical level of DBMS

- Someone has to write the DBMS software and its file manager!
- Some DB systems give the database administrator a range of physical storage options
 - Intelligent use of these options can make a very significant (and user-noticeable) difference in the way the system performs.
 - To "tune" the system properly, the DBA must understand what is happening at the physical level.
- Some of the techniques and algorithms can be used to solve other problems in other contexts

Key components of DBMS performance

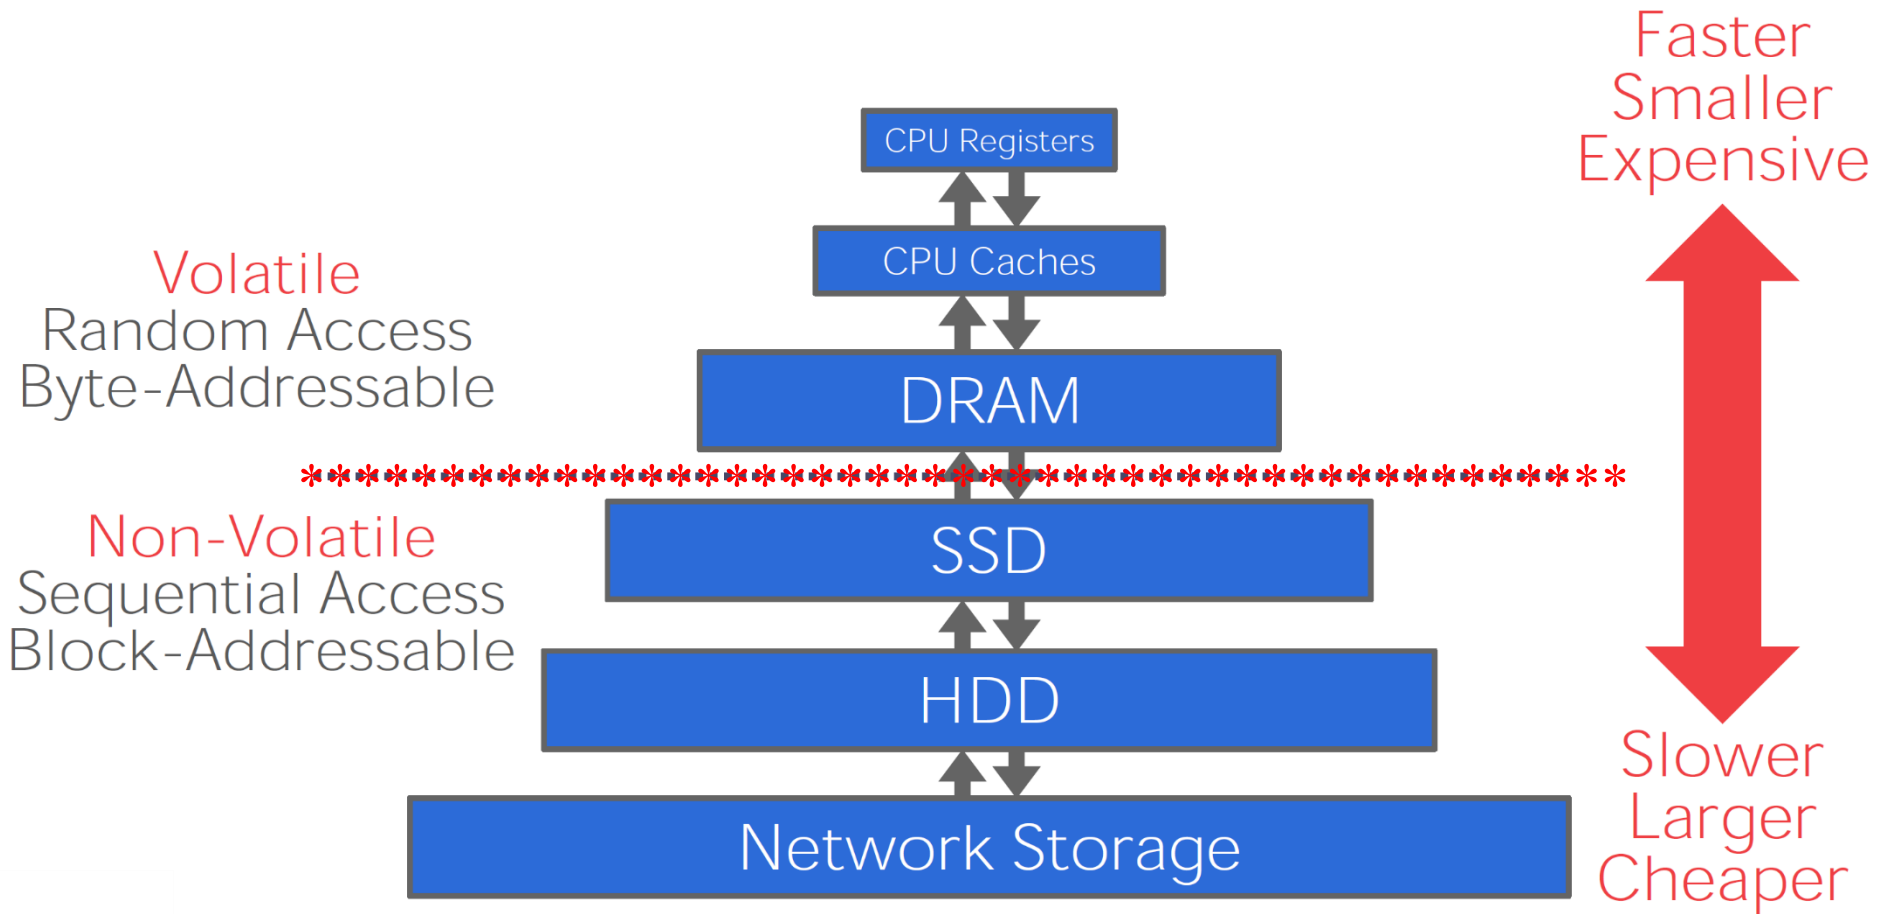
- The **performance of the DBMS Storage Engine** is often the key component of overall performance.
- There are two attributes that can be optimized:
 - 1. Response time** - defined as the time between the issuance of a command and the time that output for the command begins to be available.
 - e.g. if the command is a select statement, the time until the first row of the result appears

=> we want to minimize this

- 2. Throughput** - the number of operations that can be completed per unit time.

=> we want to maximize this

Storage Hierarchy



New trend: **Non-volatile memory** (e.g., Intel® Optane™)

<https://www.intel.com/content/www/us/en/architecture-and-technology/intel-optane-technology.html>

Storage Hierarchy - Primary storage

- **Primary storage:** Fastest media but volatile
 - Can hold subset of a database used by current transactions.
 - Volatile = data is lost when an application terminates (normally or due to a power failure or crash)
- **Cache**
 - Data and instructions in cache when needed by CPU.
 - On-board (L1) cache on same chip as CPU, L2 cache on separate chip.
 - Capacity couple of MBs, **access time few nanoseconds**
- **Main memory**
 - All active programs and data need to be in main memory.
 - Capacity couple of GBs, access time **10-100 nanoseconds**

Storage Hierarchy - Secondary & Tertiary Storage

- **Secondary storage:** non-volatile, moderately fast access time
 - Also called **online storage**
 - Stores the current version of entire database typically on a magnetic disk.
 - Access time in milliseconds
- **Tertiary storage:** non-volatile, slow access time
 - also called **off-line storage** – often used for archiving older versions of the database
 - Large capacity, access time seconds / minutes.
 - E.g. magnetic tape, optical storage

Large speed gap between Memory and Disk

- The large speed gap between Memory and Disk remains the key issue in DBMS performance.
- Time to access information in disk is the **major determining factor in system performance**.
- The *number of disk I/Os* (block accesses) is often used to measure the cost of a database operation.

Relative Daps in Access Time

0.5 ns	L1 Cache Ref	← 0.5 sec
7 ns	L2 Cache Ref	← 7 sec
100 ns	DRAM	← 100 sec
150,000 ns	SSD	← 1.7 days
10,000,000 ns	HDD	← 16.5 weeks
~30,000,000 ns	Network Storage	← 11.4 months
1,000,000,000 ns	Tape Archives	← 31.7 years

Source: <https://gist.github.com/hellerbarde/2843375>

- Each level is thousands of times faster than the level below it.
- **Dominance of I/O cost:** A modern CPU can execute millions of instructions while reading a block.

Hard Disks

- Secondary storage device of choice.
- Data is stored and retrieved in units called *disk blocks* or *pages* (typically 4 or 16 kilobytes)
- Unlike RAM, time to retrieve a disk page varies depending upon location on disk.
 - Reading several pages in sequence from a disk takes much less time than reading several random pages
- Therefore, **relative placement of pages** on disk has major impact on DBMS performance!

What's Inside A Disk Drive?





Image courtesy of Seagate Technology

Inside the HD <http://www.youtube.com/watch?v=kdmLv11n82U>



Disk Physical Structure

- Disks consist of **platters**, each with two surfaces
- Each surface consists of concentric rings called **tracks**
- Each track consists of **sectors** separated by gaps
 - Track capacities vary typically from 4 to 50 Kbytes or more
- All tracks under heads at the same time make a **cylinder** (imaginary!).
- Only one head reads/writes at any one time.




Disk Operation (Single-Platter View)

- The disk surface spins at a fixed rotational rate
- The read/write head is attached to the end of the arm and flies over the disk surface
- By **moving radially**, the arm can position the read/write head over any track




Disk Access – Service Time Components




Seek for **RED**

↑
Seek
moving arms to
position disk head
on track



Rotational latency

↑
Rotational latency
waiting for block to
rotate under head



After **RED** read

↑
Data transfer
moving data to/
from disk surface

Typically about 1% of the time is actually spent on data transfer, the rest is access time.

Performance Measures of Disks

- **Access time** – the time it takes from when a read or write request is issued to when data transfer begins. Consists of:
 - **Seek time** – time it takes to reposition the arm over the correct track.
 - 4 to 10 milliseconds on typical disks
 - **Rotational latency** – time it takes for the sector to be accessed to appear under the head.
 - 4 to 11 milliseconds on typical disks (5400 to 15000 r.p.m.)
 - **Data-transfer rate** – the rate at which data can be retrieved from or stored to the disk.
 - 25 to 100 MB per second rate, lower for inner tracks
- Access time dominated by **seek time** and **rotational latency**.
 - Disk is about 40,000 times slower than RAM

Disk speeds are dominated by access time

- For this reason, information on disk is always organized in Blocks – blocks are **basic units of transfer and storage**
 - relatively large chunks (4 or 16 kilobytes) of contiguous information that is read/written as a unit.
 - it always **reads or writes the whole block** containing a desired piece of information.
 - A system never reads or writes a single disk byte.
 - The block size B is fixed for each system.
 - Typical block sizes range from 4 to 16 kilobytes
- Mapping between logical blocks and actual (physical) sectors
 - Maintained by hardware/firmware device called disk controller.
 - Converts requests for logical blocks into (surface,track,sector) triples.

Optimization of Disk Block Access

A major goal of the design of DBMS file systems is to **minimize the time spent waiting for disk accesses**. 3 ways this is done:

- 1) **Store related information on the same or nearby blocks:**
read and write of data on ***contiguous disk blocks*** and eliminates seek time and rotational delay time for all but the first block transfer
 - Files may get fragmented over time (if data is inserted to/deleted from the file) => reorganize the database files to speed up access
- 2) **Keeping copies of recently-used information in buffers in memory**, so that if the same information is needed again it can be accessed without having to go to the disk again
- 3) **Parallelism** - spreading information across multiple disks, so that several disks can be going through the physical operations needed to access information at the same time

Example: reading two disk blocks

- Assume
 - average seek time = 10 ms
 - average rotational latency = 3 ms
 - transfer time for 1 block = 0.01875 ms
- **Adjacent** block on same track
 - access time = $10 + 3 + 2 \times (0.01875)$ ms = 13.0375 ms
- **Random** block
 - access time = $2 \times (10 + 3 + 0.01875)$ ms = 26.0375 ms

RAID Technology

Parallelizing Disk Access using RAID Technology

- **RAID: Redundant Arrays of Independent Disks**
 - an array of independent disks **acting as a single higher-performance logical disk**, providing:
 - **high capacity** and **high speed** by using multiple disks in parallel
 - reduce the large speed gap between disks and the memory
 - **high reliability** by storing data redundantly, so that data can be recovered even if a disk fails

RAID goals




RAID systems seek :

- **to improve throughput** by a technique known as **striping**, in which a single file is spread over multiple disks.
 - **Parallelize** large accesses to reduce response time: multiple accesses to different parts of the same file can often be performed in parallel (assuming that the parts being accessed are on different disks).
- **to improve reliability** by **replication** of data, so that if a disk fails, the data it contained is available somewhere else.
 - => improve throughput for reads -if there are multiple copies of an item, then any copy can be read.
 - but creates an issue on write though - since all copies must be updated.

RAID 0 - a.k.a. Striping

- Requires two or more disks
- No lost drive space due to striping
- Fastest read and write performance.
- Raid level 0 has no redundant data and hence has the best write performance at the risk of data loss
 - Offers no data protection.
- The more disks, the more risk.

Used in high-performance applications where data loss is not critical



Data striping

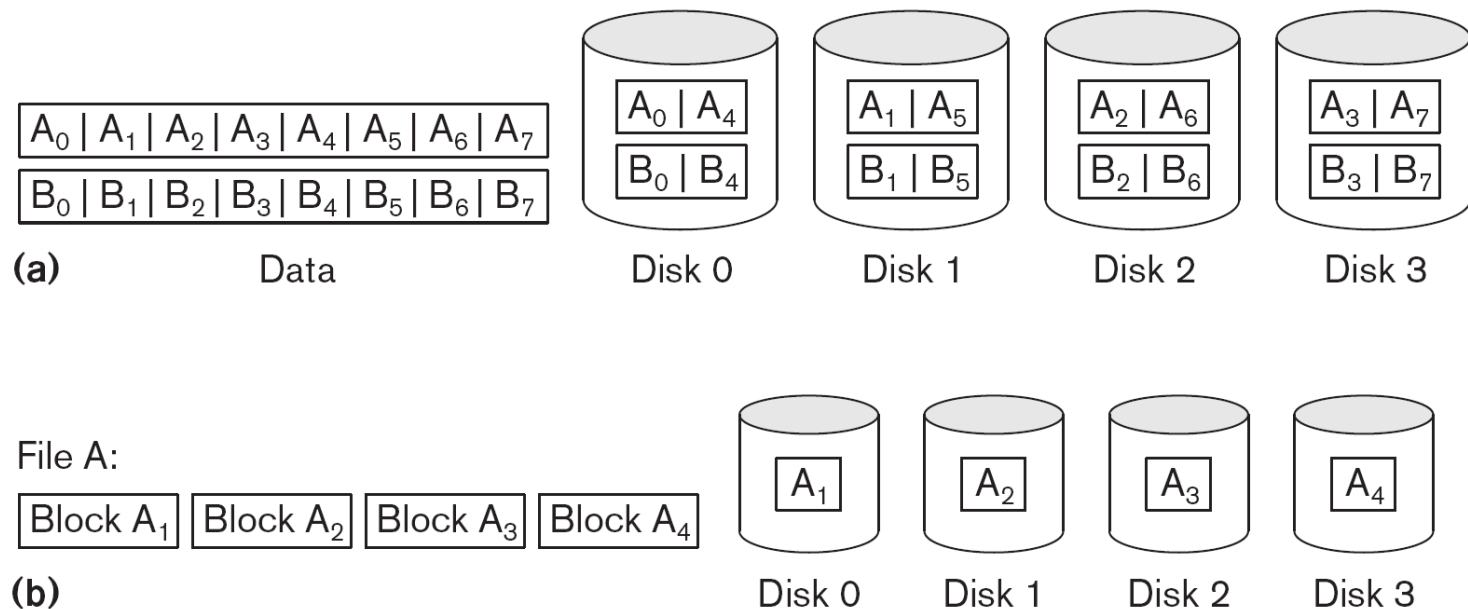
- RAID uses a concept called **data striping** = distribute blocks over n disks in a round robin fashion.
 - Make disks appear as a single large, fast disk.
 - Requests for different blocks can run in parallel if the blocks reside on different disks

Figure 17.13

Striping of data across multiple disks.

(a) Bit-level striping across four disks.


(b) Block-level striping across four disks.



RAID 1 - a.k.a. Mirroring


- Raid level 1 uses mirrored disks
- Write speed of one disk
- Read speed of two disks
- Capacity is equal to the size of one

Popular for applications such as storing log files in a database system.




RAID 5 - Striping with Distributed Parity

- Considered best compromise between speed and storage efficiency:
 - Good performance (as blocks are striped) but slower writes due to parity
 - Good redundancy (distributed parity)
- Requires 3 or more drives
- Stripe across all drives with **parity**
- Can loose 1 drive and still function
- Capacity is **$n-1$** where **n** is number of drives in array

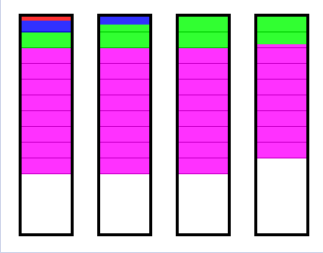
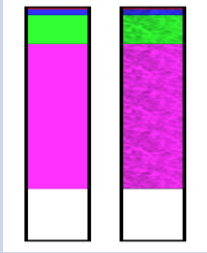
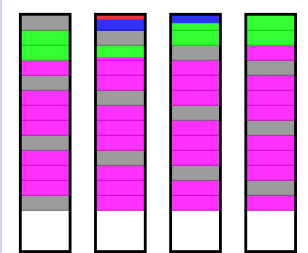
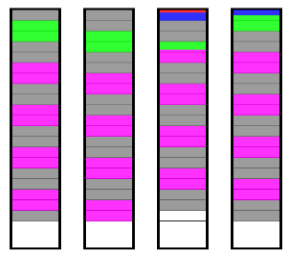


RAID 6 - RAID 5 on Steroids

- 4 or more disk
- Is a stripe with two parity drives
- Can loose two drives and still function
- Capacity is $n-2$ where n is number of drives in array
- Protect against up to two disk failures by using just two redundant disks



Comparison of Single RAID Levels

	RAID 0	RAID 1	RAID 5	RAID 6
Diagram				
Description	Striping	Mirroring	Striping with Parity	Striping with Dual Parity
Minimum Disks	2	2	3	4
Array Capacity	No. of Drives x Drive Capacity	Drive Capacity	(No. of Drives - 1) x Drive Capacity	(No. of Drives - 2) x Drive Capacity

Comparison of Single RAID Levels

	RAID 0	RAID 1	RAID 5	RAID 6
Storage Efficiency	100%	50%	(Num of drives – 1) / Num of drives	(Num of drives – 2) / Num of drives
Fault Tolerance	None	1 Drive failure	1 Drive failure	2 Drive failures
High Availability	None	Good	Good	Very Good
Degradation during <u>rebuild</u>	NA	<ul style="list-style-type: none"> • Slight degradation • Rebuilds very fast 	<ul style="list-style-type: none"> • High degradation • Slow Rebuild (due to write penalty of parity) 	<ul style="list-style-type: none"> • Very High degradation • Very Slow Rebuild (due to write penalty of dual parity)

Understanding the Parity


- RAID 5 and RAID 6 store parity information against data for rebuild
- Parity can be calculated using a simple XOR
- eg– “ABCDEFGHijkl” on a 4 disk RAID 5 array

Disk 1	Disk 2	Disk 3	Disk 4
A (01000001)	B (01000010)	C (01000011)	{P – 01000000}
Parity {P}	D	E	F
G	Parity {P}	H	I
J	K	Parity {P}	L

- If Disk 2 fails then the data “B” can be recalculated as
(01000001 XOR 01000011 XOR 01000000) => 01000010 => B
- More info @ <http://www.youtube.com/watch?v=LTq4pGZtzh0>

RAID 10 a.k.a. 1+0

Diagram



Description

Mirroring then Striping

Minimum Disks

Even number > 4

Maximum Disks

Controller Dependant

Array Capacity

$(\text{Size of Drive}) * (\text{Number of Drives}) / 2$

Storage Efficiency

50%

Fault Tolerance

Multiple drive failure as long as 2 drives from same RAID 1 set do not fail


High Availability

Excellent

Choice of RAID Level

- Factors in choosing RAID level
 - Monetary cost
 - Performance: Number of I/O operations per second, and bandwidth during normal operation
 - Performance during failure
 - Performance during rebuild of failed disk
- RAID 0 is used only when data safety is not important
 - E.g. data can be recovered quickly from other sources
- Level 6 is rarely used since levels 1 and 5 offer adequate safety for most applications


Networked Storage



Source: <http://www.youtube.com/watch?v=2T99tW1KEMc>

Storage Area Network (SAN)


- Online storage peripherals are configured as nodes on a **high-speed network** and can be attached and detached from servers in a very flexible manner
- Servers see SAN as a virtual drives
- **Dedicated access** - each part of the SAN is dedicated to each server
- **Block based storage**



- Storage traffic segregating from the rest of the LAN. It typically uses Fiber Channel connectivity

Network Attached Storage (NAS)

- File Server optimized to serve files over the main LAN (OS dedicated to file system)
- File based storage
- Servers see NAS as a Network Share (need to map it to a drive)
- Suitable for sharing files



Database File Organization


Database Storage

Database Storage addresses 2 problems:

- **Problem #1:**How the DBMS represents the database in files on disk.
 - File Storage
 - Page Layout
 - Record Layout
- **Problem #2:**How the DBMS manages its memory and efficiently move data back-and-forth from disk.

File Organization

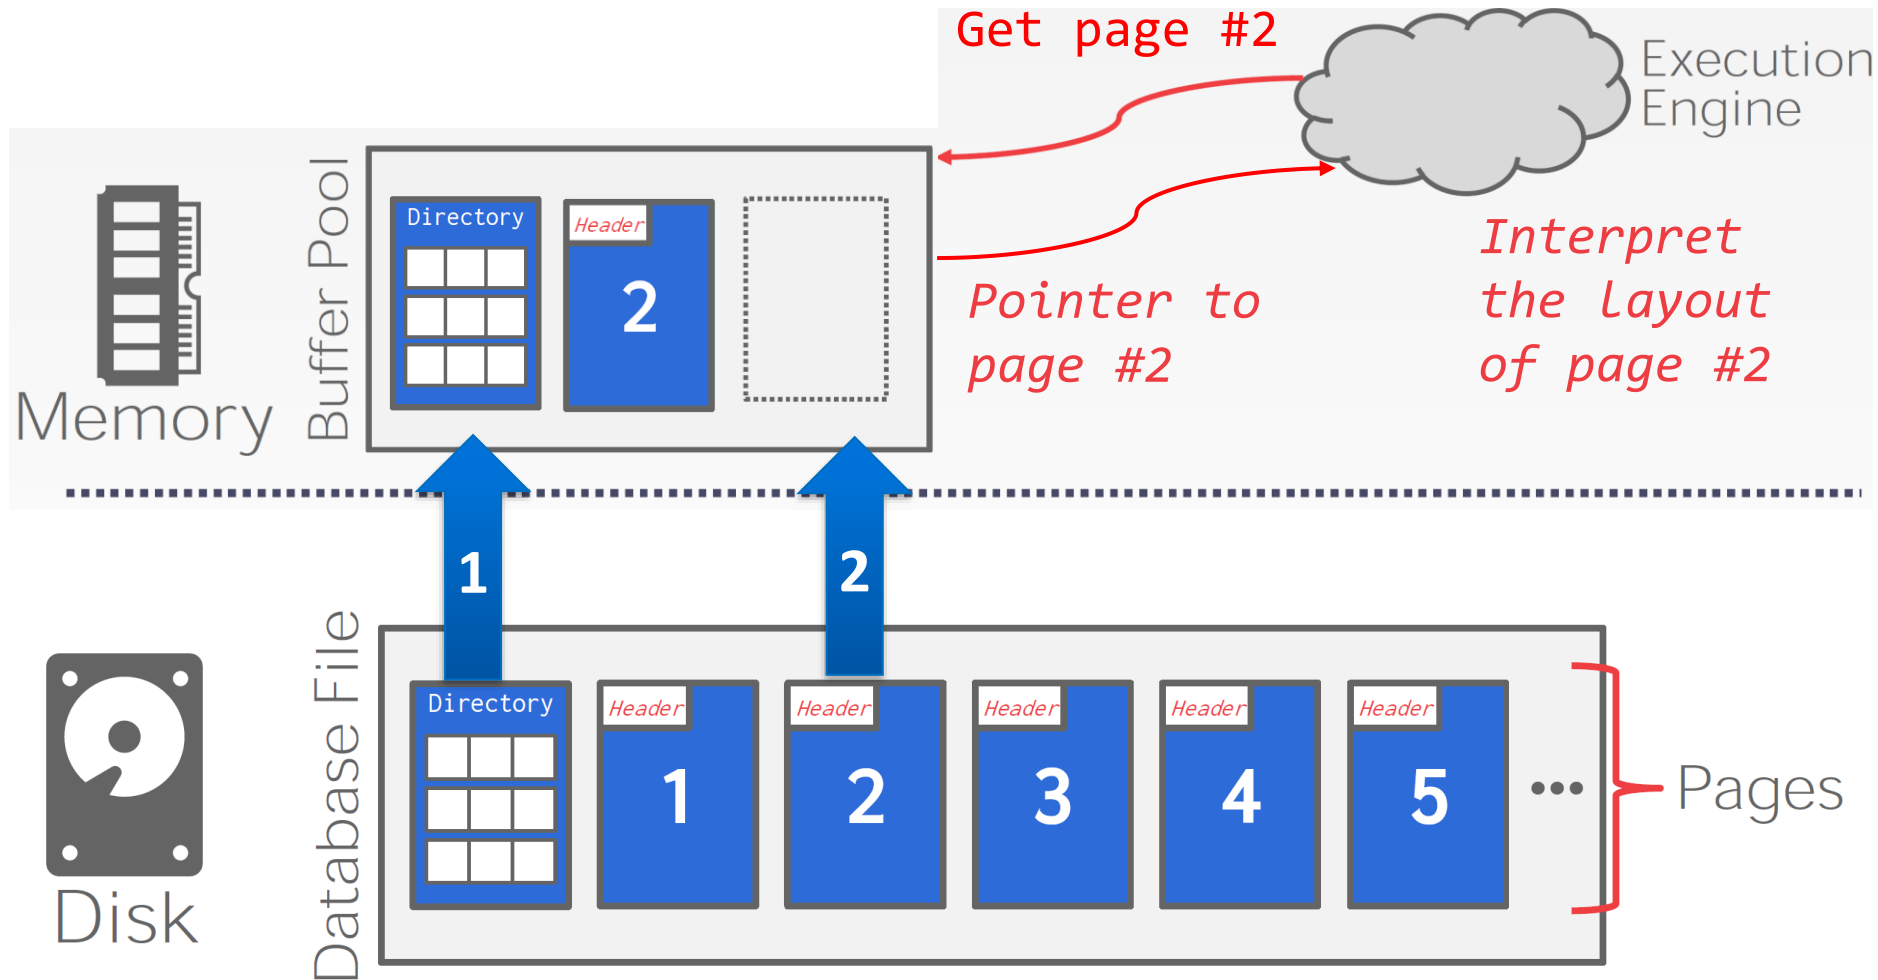
- The DBMS stores a database as one or more *files* on disk.
 - The OS doesn't know anything about the contents of these files.
- Each file has a collection of *pages* (aka blocks)
- Each *page* is a sequence of *records*.
- A record is a sequence of *fields*.



Storage Manager

- The **storage manager** is responsible for maintaining a database's files.
 - It organizes the files as a collection of pages.
 - Tracks data read/written to pages.
 - Tracks the available space.

How the Storage Engine Works?









Why NOT use the OS?

- In DBMS the OS is NOT used for moving data pages in and out of memory.
- DBMS (almost) always wants to control things itself and can do a better job at it.
 - Flushing dirty pages to disk in the correct order.
 - Specialized prefetching.
 - Buffer replacement policy.
 - Thread/process scheduling.

Database Pages

- A **page** is a **fixed-size** block of data (4 to 16KB).
 - It can contain records, meta-data, indexes, log records...
- Each page is given a unique identifier.
- The DBMS uses an indirection layer (i.e., a **dictionary**) to map page ids to physical locations.

4KB	 SQLite
	 IBM DB2
	 ORACLE®
<hr/>	
8KB	 Microsoft® SQL Server®
	 PostgreSQL
<hr/>	
16KB	 MySQL™

Heap File

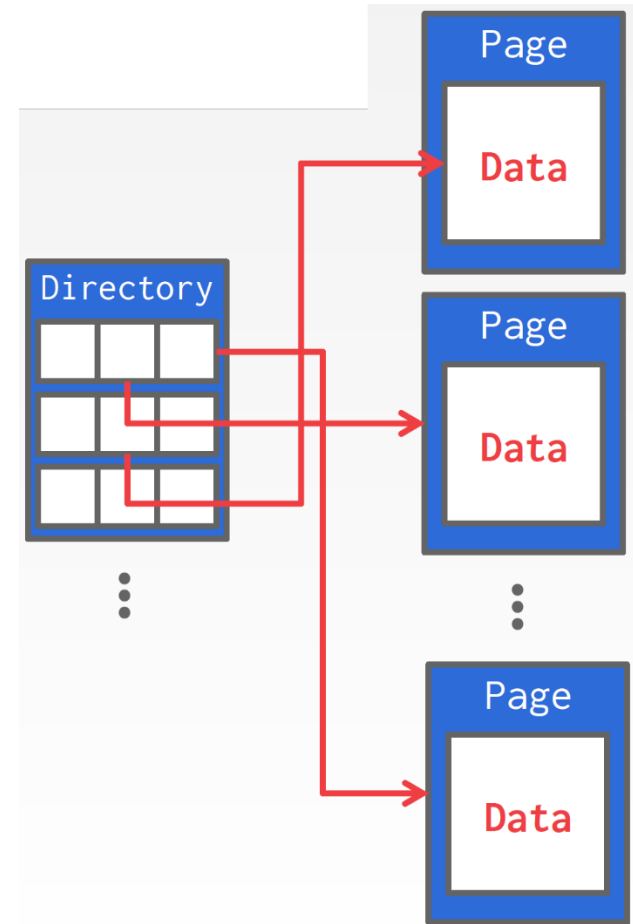
- DBMSs manage pages in files on disk in different ways:
 - Heap File Organization
 - Sequential / Sorted File Organization
 - Hashing File Organization
- A **heap** file is an unordered collection of pages where records are stored in random order. Supported operations include:
 - Create / Get / Write / Delete Page
 - Iterate over all pages
- Need meta-data to keep track of what pages exist and which ones have free space. Typically using a **Page Directory**

Heap File

- A **linear search** through the file records is necessary to search for a record
 - This requires reading and searching half the file blocks on the average, and is hence quite expensive
- Record insertion is efficient as new records are inserted at the end of the file
- Reading the records in order of a particular field requires sorting the file records.
- Deleted rows create gaps in file
 - File must be periodically compacted to recover space


Heap File – Page Directory

- The DBMS maintains special pages that **tracks the location of data pages** in the database files.
- The directory also records the **number of free slots** per page.
- The DBMS has to make sure that the directory pages are in sync with the data pages.




Page Layout

- Every page contains a **header** of meta-data about the page's contents:
 - Page Size
 - Checksum
 - DBMS Version
 - Compression Information
- The page data is organized using 3 approaches:
 - Row-oriented
 - Column-oriented
 - Log-structured



Row-oriented Storage

- A data layout that contiguously stores the values belonging to the columns that make up the entire row.
- The most common layout scheme is called **slotted pages**.
- The slot array maps "slots" to the records' **starting position** offsets.
- The header keeps track of:
 - The # of used slots
 - The offset of the starting location of the last slot used.




Row-oriented Storage

- Row storage model (aka "n-ary storage model")
- The DBMS stores all attributes for a single record contiguously in a page.
- Ideal for On-line Transaction Processing (OLTP) workloads where queries tend to operate only on an individual entity and insert heavy workloads.
- Advantages
 - Fast inserts, updates, and deletes.
 - Good for queries that need the entire tuple.
- Disadvantages
 - Not good for scanning large portions of the table and/or a subset of the attributes.

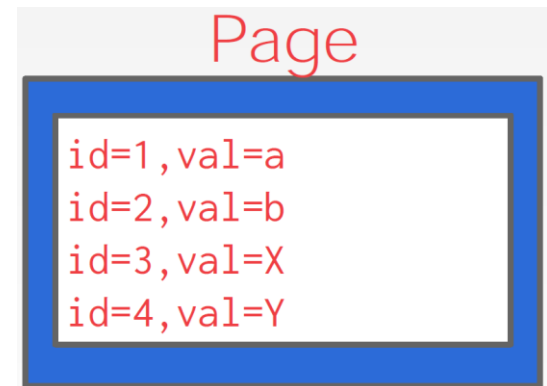
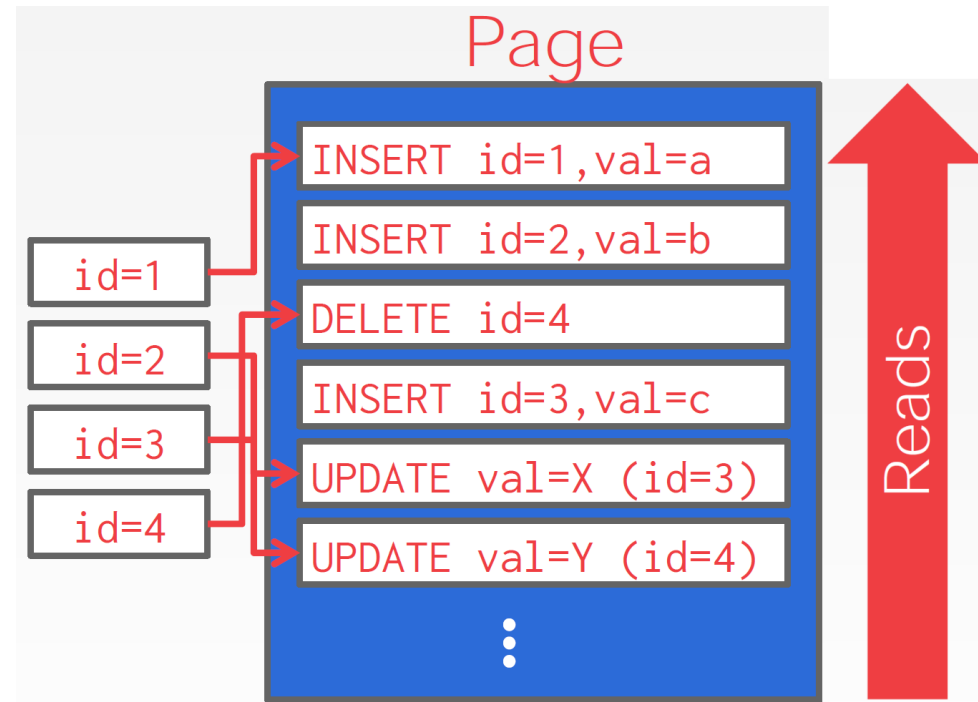
Log-structured Organization

- Instead of storing records in pages, the DBMS only stores **log records**.
- The system appends log records to the file of how the database was modified:
 - Inserts store the entire tuple.
 - Deletes mark the tuple as deleted.
 - Updates contain the **delta** of just the attributes that were modified.



Log-structured Organization

- To read a record, the DBMS scans the log **backwards** and "recreates" the record
- Indexes can be used to allow it to jump to locations in the log.
- Periodically **compact** the log to improve read performance.



Record Layout

- The DBMS assigns each record a **unique record identifier** to keep track of individual records, commonly:

Record id = $\text{pageId} + \text{slot\#}$

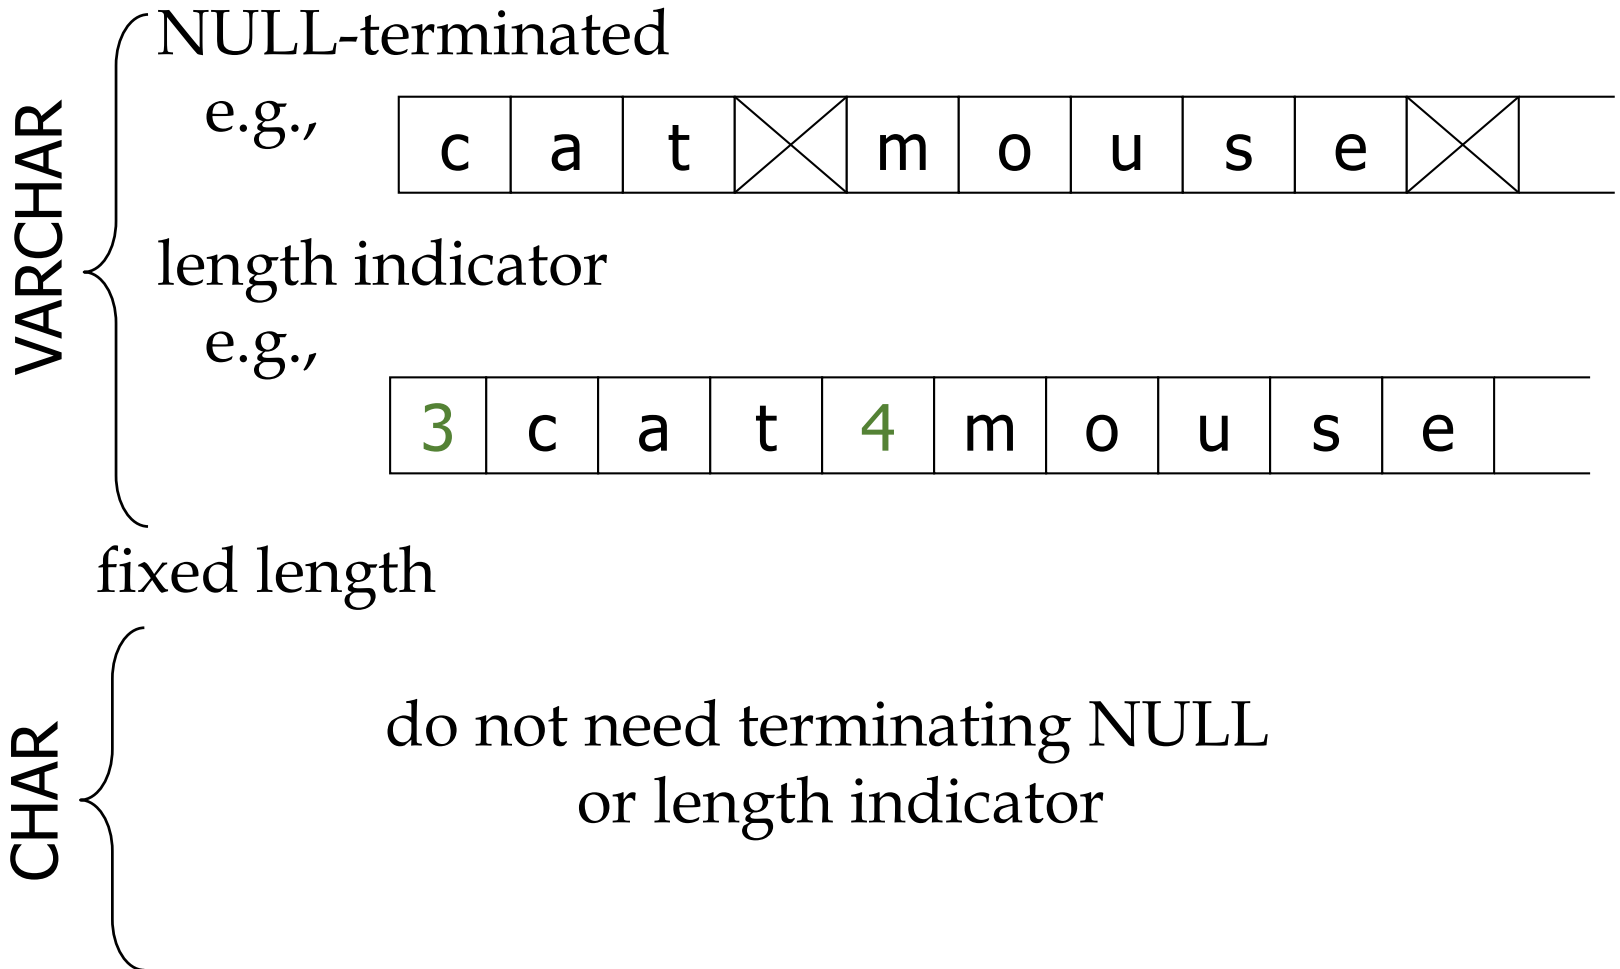
- A record is essentially a sequence of bytes.
- It is the job of the DBMS to interpret those bytes into attribute **types** and **values**.
- DBMS's **catalog** contain the schema information about tables that the system uses to figure out the record's layout.

Record

<i>Header</i>	a	b	c	d	e
---------------	---	---	---	---	---


```
CREATE TABLE foo (  
  a INT PRIMARY KEY,  
  b INT NOT NULL,  
  c INT,  
  d DOUBLE,  
  e FLOAT  
);
```

Handling String of characters



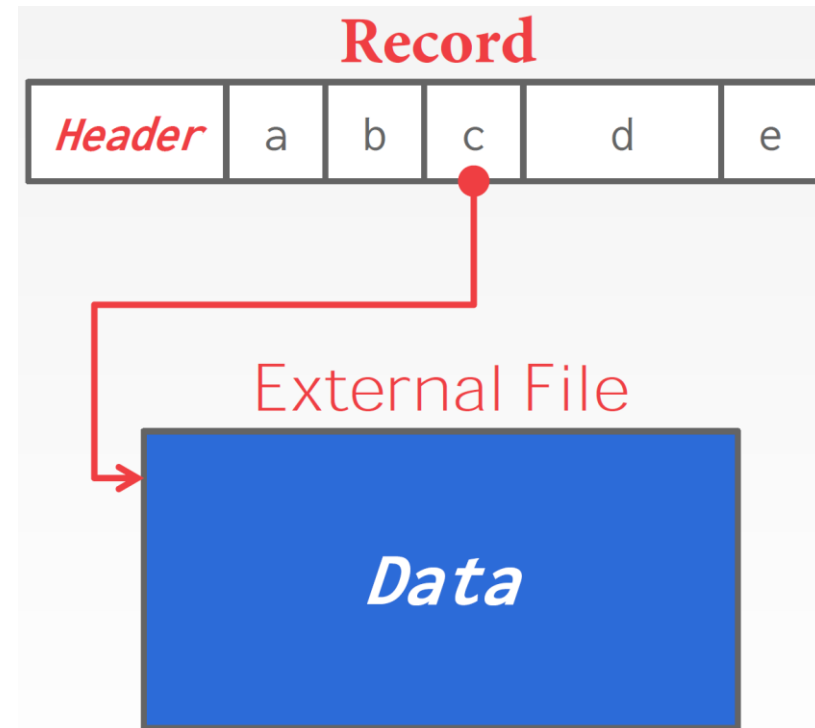
Storing Large Values

- Most DBMSs don't allow a record to exceed the size of a single page.
- To store values that are larger than a page, the DBMS uses separate **overflow storage pages**.



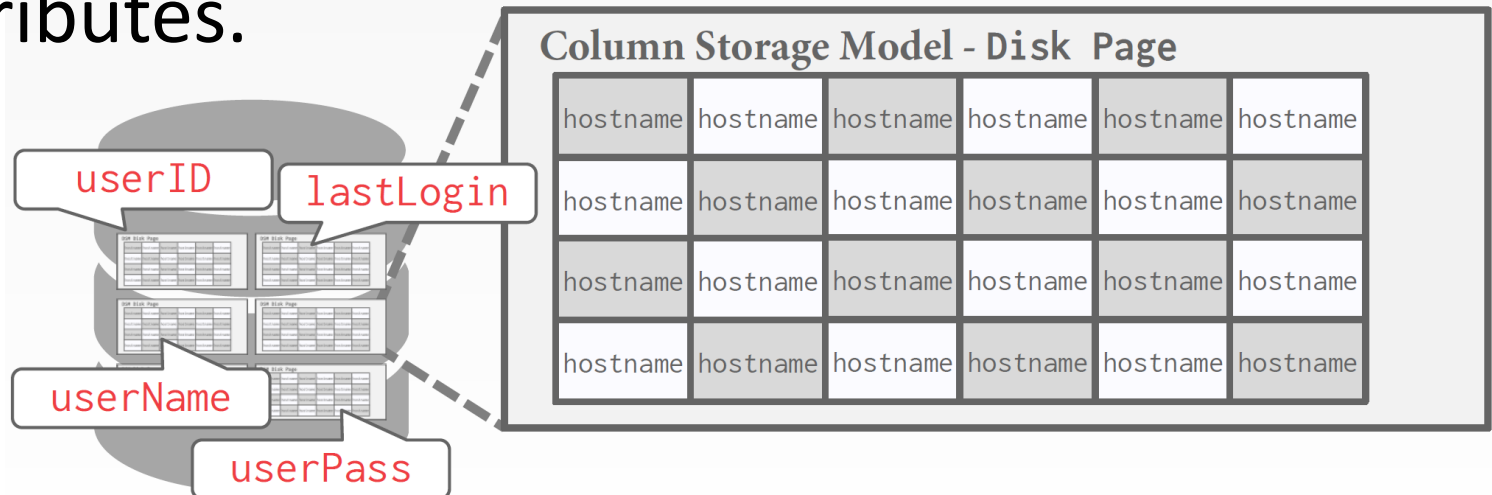
External Value Storage

- Some systems allow you to store a really large value in an external file. Treated as a **BLOB** type.
 - Oracle: **BFILE** data type
 - Microsoft: **FILESTREAM** data type
- The DBMS **cannot** manipulate the contents of an external file
 - No durability protections
 - No transaction protections



Column Storage Model

- The DBMS stores the values of a single column for all records contiguously in a page.
- Ideal for On-line Analytical Processing (OLAP) workloads where read-only queries perform large scans over a subset of the table's attributes.



Column Storage Model

- **Advantages**

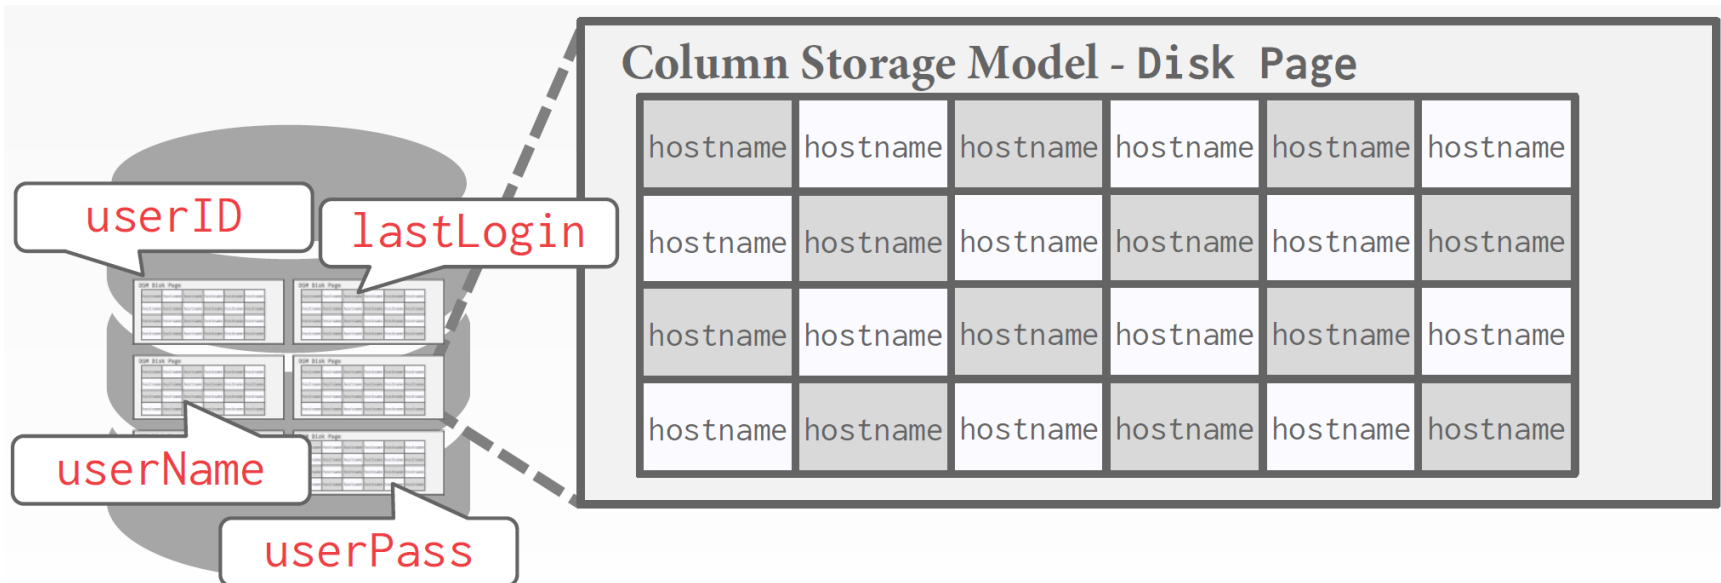
- Reduces the amount wasted I/O because the DBMS only reads the data that it needs.
- Better performing aggregation over large volumes of data. Or for queries that only need a few columns from a wide table.
- Better query processing and data compression

- **Disadvantages**

- Slow for point queries returning many columns because of record stitching
- Slow inserts, updates, and deletes because of record splitting

Example

```
SELECT COUNT(U.lastLogin),  
       EXTRACT(month FROM U.lastLogin) AS month  
FROM useracct AS U  
WHERE U.hostname LIKE '%.gov'  
GROUP BY EXTRACT(month FROM U.lastLogin)
```



Getting Columns Belonging to a Record

To put the records back together we can use:

- **Choice #1: Fixed-length Offsets**
 - Each value is the same length for an attribute. Then when the system wants the attribute for a specific record, it knows how to jump to that spot in the file.
- **Choice #2: Embedded Tuple Ids**
 - Each value is stored with its record id in a column.

Most DBMSs use fixed-length offsets

Offsets

	A	B	C	D
0				
1				
2				
3				

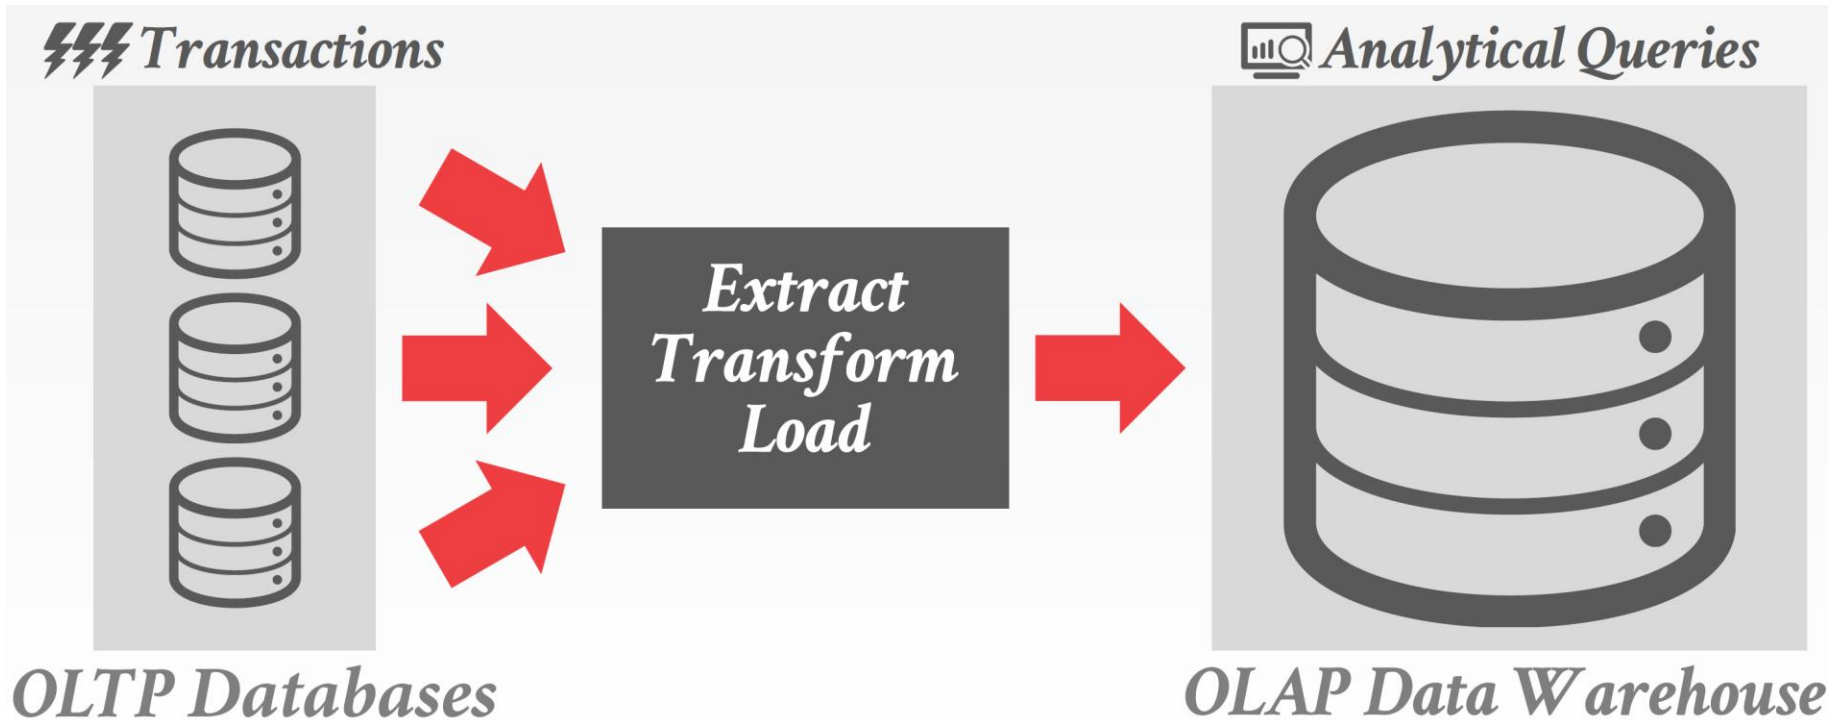
Embedded Ids

	A	B	C	D
0				
1				
2				
3				

OLTP vs. OLAP

- On-line Transaction Processing:
 - Simple queries that read/update a small amount of data that is related to a single/few entities in the database.
- On-line Analytical Processing:
 - Complex read intensive queries that read a lot of data to compute aggregates.
- The DBMS can store records in different ways that are better for either OLTP or OLAP workloads:
 - OLTP = Row Store
 - OLAP = Column Store


Extract Transform Load (ETL)



Buffer Manager

Buffer Pool

- Data must be in RAM for DBMS to operate on it!
- Buffer pool is an in-memory **cache** of pages read from disk
- Buffer Pool = Memory region organized as an array of fixed-size pages.
 - An array entry is called a **frame**
- When the DBMS requests a page, a copy is placed from disk into one of these frames
- Enables the higher level DBMS components to assume that the needed data is in main memory




Buffer Pool Manager

Why not use the Operating System for the task??

- DBMS may be able to anticipate access patterns
- => Hence, may also be able to perform prefetching
- DBMS needs the ability to force pages to disk

Buffer Pool Metadata

- The **page table** keeps track of database pages that are currently in the buffer pool frames.
- Also maintains additional meta-data per page:
 - **Pin Counter**
 - **Dirty Flag**
- Page table contains:



<frame#, pageid, pin_count, dirty>

Buffer Pool Metadata

- **Pin Counter:** Tracks the number of threads that are currently accessing the frame.
 - A thread has to increment the counter before they access the frame.
 - If Pin Counter > 0 , then the storage manager is not allowed to evict that frame from memory.
- **Dirty Flag:** set to 1 when a page is modified
 - This indicates to storage manager that the page must be written back to disk

When a Page is Requested ...

- If requested page is not in pool:
 - If no free frame available, choose a frame for *replacement*
Only frames with `pin_count == 0` are candidates!
 - If frame is not *dirty*, then simply evict it.
 - If frame is *dirty*, write it to disk to ensure that its changes are persisted.
 - Read requested page into chosen frame
- *Pin* the page and return its address

More on Buffer Management

- Requestor of page must eventually unpin it (to indicate it is no longer needed) + indicate whether page has been modified: *dirty* bit is used for this.
- Page in pool may be requested many times,
 - A *pin count* is used.
 - To pin a page, `pin_count++`
 - A page is a candidate for replacement if *pin count* == 0 (“unpinned”)

Buffer Replacement Policies

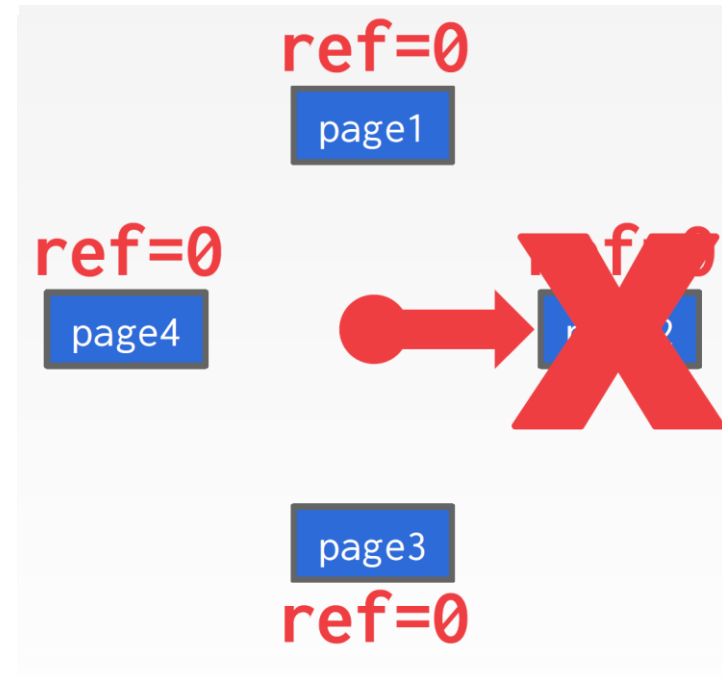
- The buffer pool provides space for a limited number of disk pages
- When the DBMS needs to free up a frame to make room for a new page, it must decide which page to *evict* from the buffer pool
- Frame is chosen for replacement by a *replacement policy*:
 - Least-recently-used (LRU)
 - LRU-K
 - Private Pool Space per Query
 - Priority hints

Least Recently Used (LRU) Policy

- Least Recently Used (LRU) Policy:
 - Maintain a timestamp of when each page was last accessed.
 - When the DBMS needs to evict a page, select the one with the oldest timestamp.

Clock

- Approximation of LRU without needing a separate timestamp per page.
 - Each page has a reference bit
 - When a page is accessed, set to 1
- Organize the pages in a circular buffer with a "*clock hand*":
 - Upon sweeping, check if a page's bit is set to 1.
 - If yes, set to zero. If no, then evict
 - Clock hand remembers position between evictions



LRU Problems

- LRU and CLOCK replacement policies are susceptible to ***sequential flooding***
 - A query performs a sequential scan that reads every page => causing trashing the buffer pool contents due to a sequential scan
 - This **pollutes** the buffer pool with pages that are read once and then never again
- The most recently used page is actually the most unneeded page.

Sequential Flooding

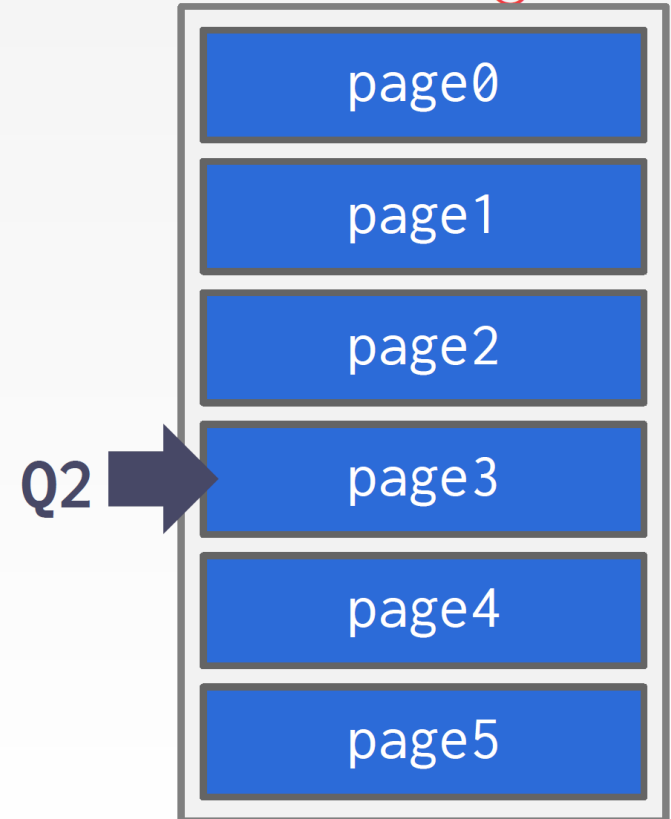
Q1 `SELECT * FROM A WHERE id = 1`

Q2 `SELECT AVG(val) FROM A`

Buffer Pool



Disk Pages



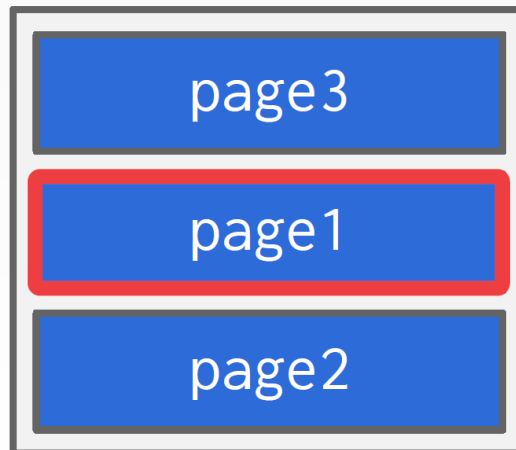
Sequential Flooding

Q1 `SELECT * FROM A WHERE id = 1`

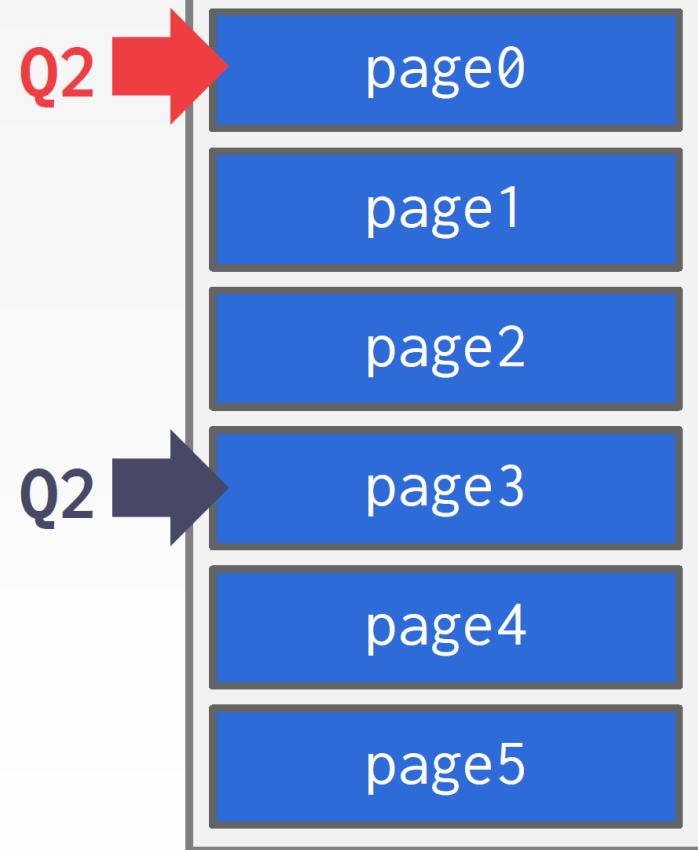
Q2 `SELECT AVG(val) FROM A`

Q3 `SELECT * FROM A WHERE id = 1`

Buffer Pool



Disk Pages



LRU-K

- Track the history of the last **K** references as timestamps and compute the interval between subsequent accesses
- The DBMS then uses this history to estimate the next time that page is going to be accessed

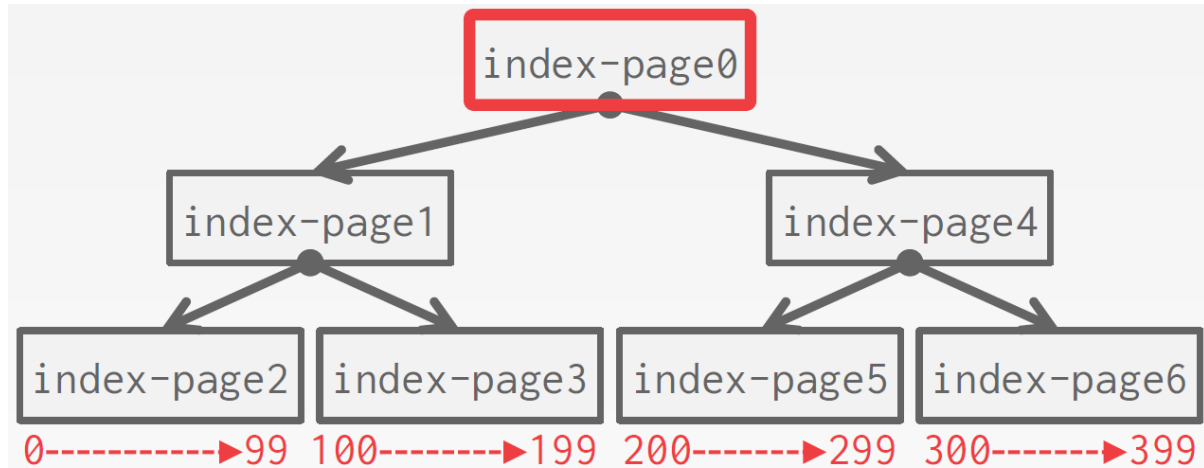
Private Pool Space per Query

- The DBMS chooses which pages to evict on a per transaction/query basis. This minimizes the pollution of the buffer pool from each query.
 - Keep track of the pages that a query has accessed.

Priority hints

- The DBMS knows what the context of each page during query execution
- It can provide hints to the buffer pool on whether a page is important or not

e.g., The **root** of a B+Tree index is always kept on the Buffer Pool

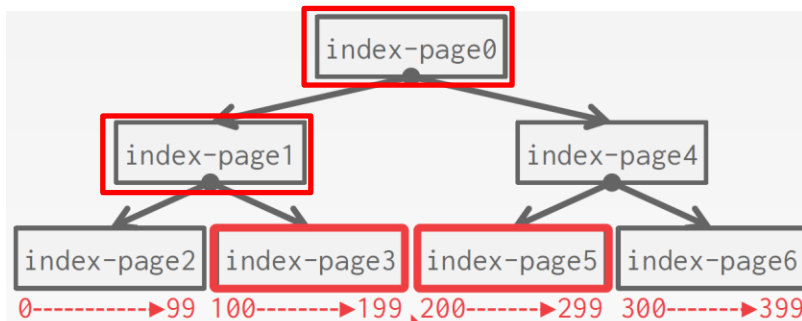


Pre-fetching


- The DBMS can prefetch pages based on the query plan
 - Sequential Scans
 - Index Scans

Q1

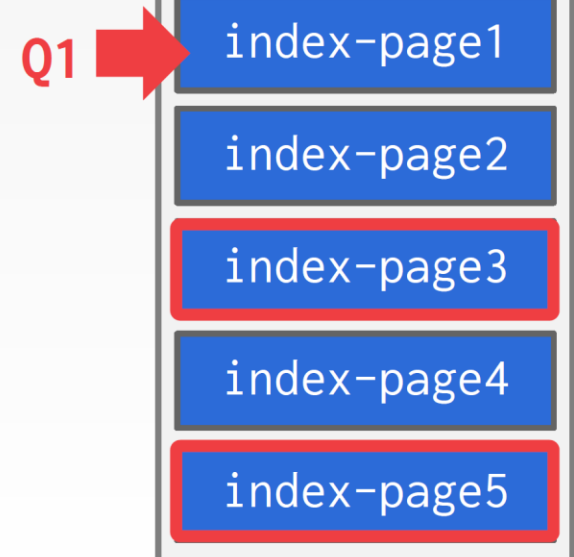
```
SELECT * FROM A
WHERE val BETWEEN 100 AND 250
```



Buffer Pool



Disk Pages



Summary

- I/O times dominate DBMS performance
- Techniques to improve I/O time:
 - Buffer Pool manager: leverages the semantics about the query plan to make better Pre-fetching and Eviction decisions.
 - RAID Technology
 - Store related data on contiguous blocks
 - To keep good performance, the DBMS must occasionally rebuild the database files to merge in the overflow pages and reclaim unused blocks