

Hashing

CMPT 6o6




Read Chapter 16
(Hashing Sections)

Outline

- ① Extendible Hashing
- ② Linear Hashing

Hash Tables in Main Memory

- In a hash table a hash function maps search key values to array elements
 - The array can either contain the data objects, or
 - Linked lists containing data objects (often called *buckets*)
- Hash functions generate a value between 0 and $B-1$
 - Where B is the number of buckets
 - A record with **search key K** is **stored in bucket $h(K)$**

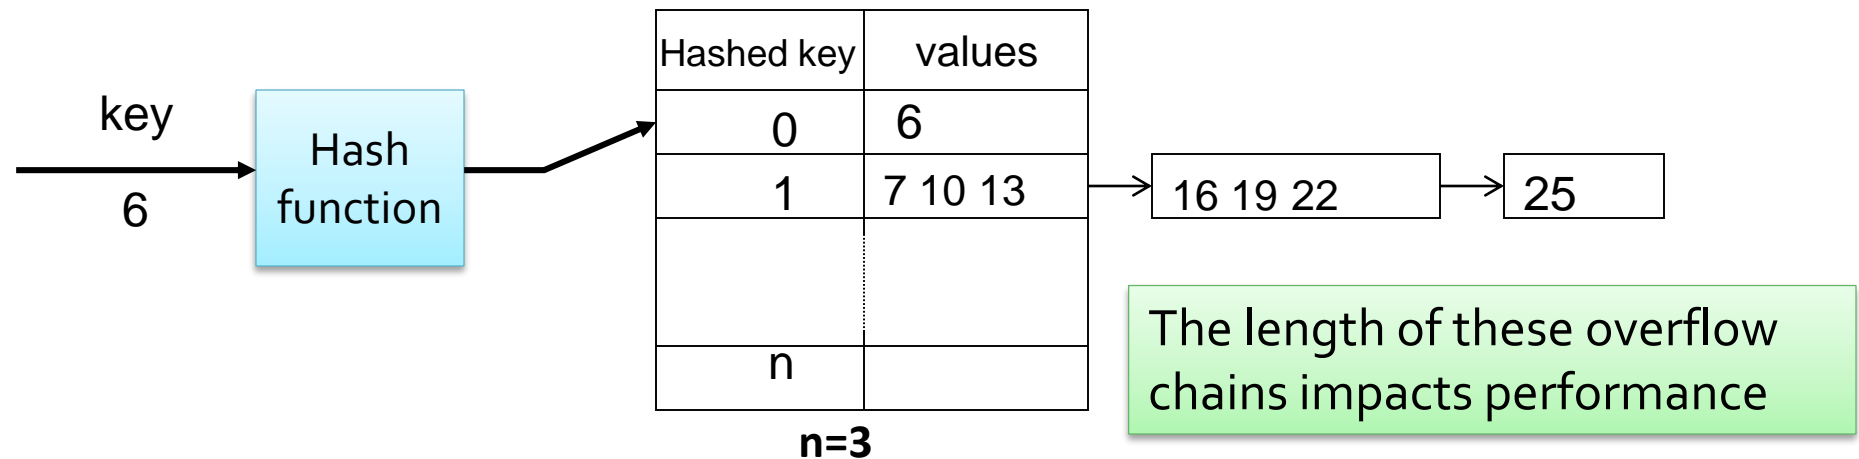


- A typical hash function is a bit representation of (the Search Key Value **modulus** the Number of Buckets)
- Hash indexes do not support range lookup

Hash Index



- A hash index stores **key-value pairs** based on a *hash function* => *an array of pointers to buckets*
- Fast data structures that support **$O(1)$** look-ups
- The hash function often is **$K \bmod B$**
 - Where **K** is the key value and **B** is the number of buckets
- Collision is resolved by placing new values in an **overflow bucket**



Hash Functions

- Hash (any input key) return an **integer representation** of that key. DBMS hash tables use a fast hash function with a low collision rate:
- [CRC-64](#) (1975)
 - Used in networking for error detection
- [MurmurHash](#) (2008)
 - Designed to a fast, general purpose hash function
- [Google CityHash](#) (2011)
 - Designed to be faster for short keys (<64 bytes)
- [Facebook XXHash](#) (2012)
 - It is proposed in two flavors, 32 and 64 bits
- [Google FarmHash](#) (2014)
 - Newer version of CityHash with better collision rates



Static Hashing vs. Dynamic Hashing

- In static hashing the number of buckets is fixed (assumes you know the number of entries ahead of time)
 - If the file grows some buckets will have overflow chains, reducing efficiency

=> There should be enough buckets so that there are few overflow blocks
- Two versions of dynamic hashing that allow **dynamic File Expansion**
 - Extensible hashing
 - Linear hashing

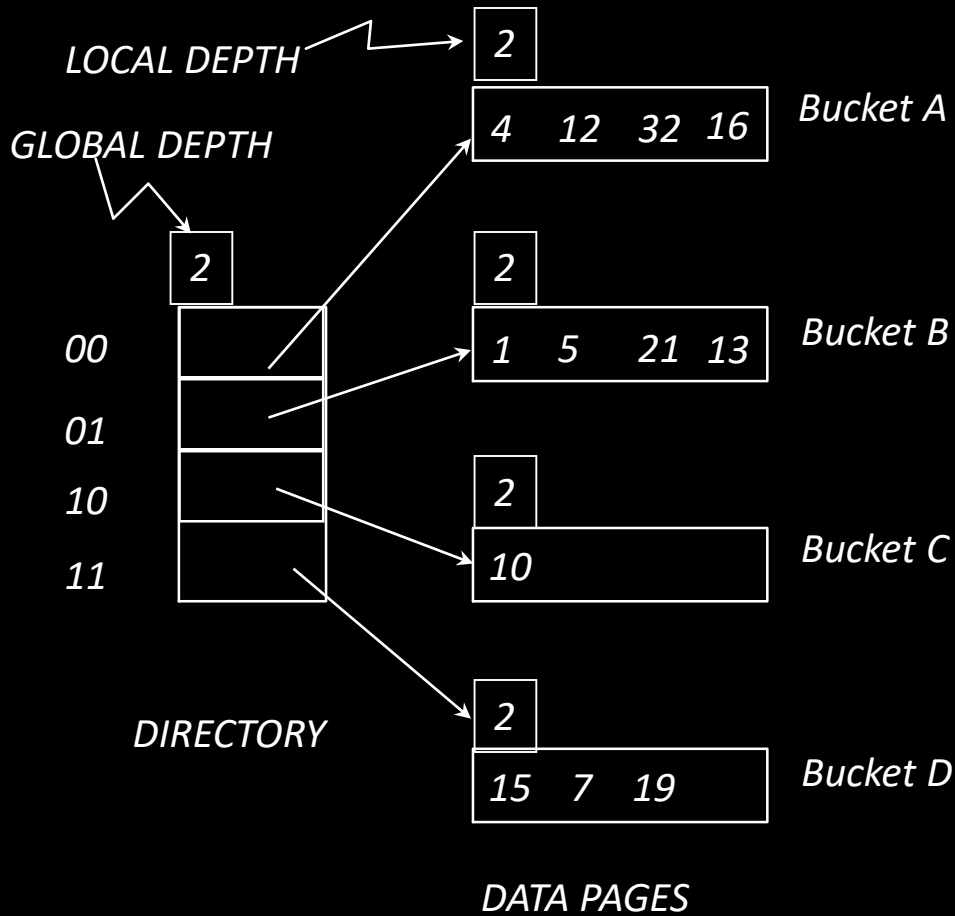
Extendible hashing

Extendible Hashing Structure

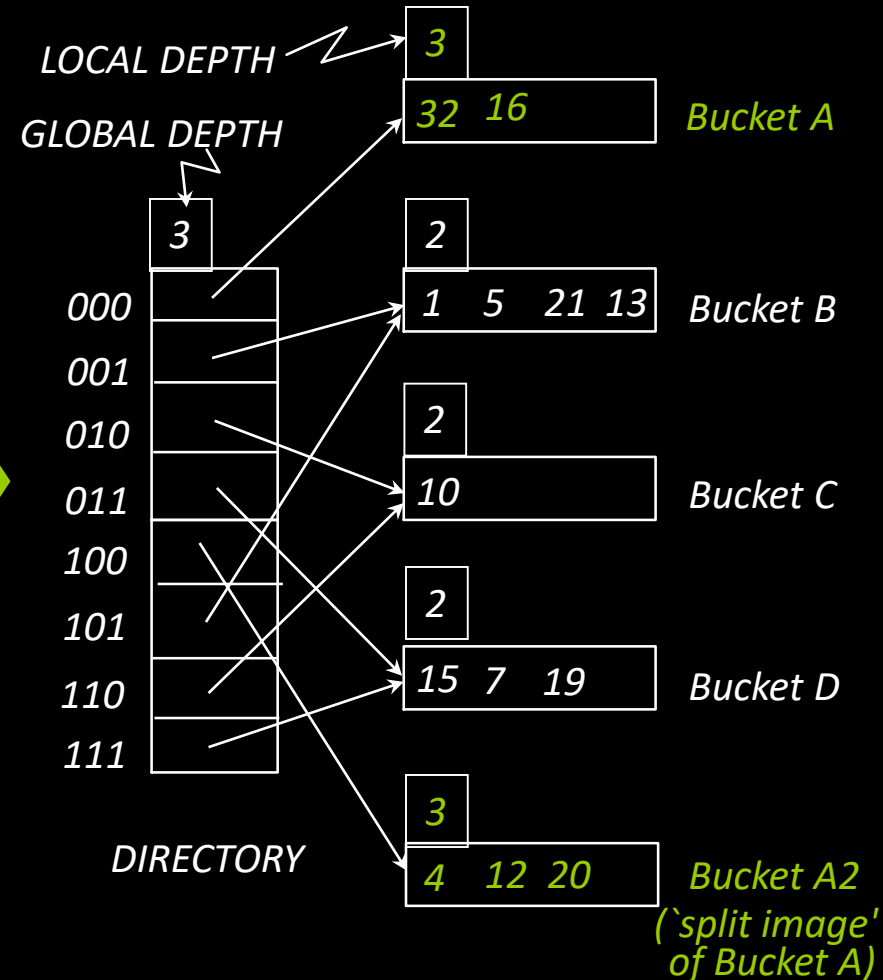
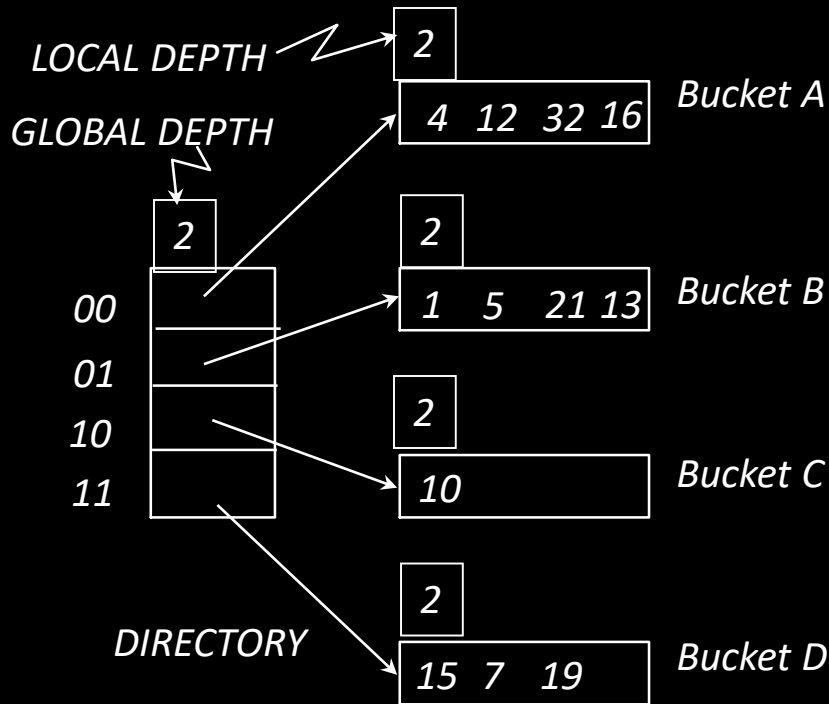
- In extendible hashing there is a **directory** to buckets
 - The directory is an array of pointers to buckets
 - As the array only contains pointers it is relatively small, so it can usually fit in memory
 - The directory size is doubled when overflow occurs
 - New buckets are only created as necessary
- The hash function computes a sequence of bits
 - The directory (and the associated buckets) uses the last **i bits** of the hash value
 - The directory will have **2^i** entries
 - When the directory grows, **$i+1$** bits of the hash value are used

Extendible Hashing Example

- Directory is array of size 4.
- Assume a bucket capacity of 4
- What info this hash data structure maintains?
 - **Global Depth:** # of bits needed to tell which bucket an entry belongs to
 - **Local Depth:** # of bits used to determine if an entry belongs to a specific bucket



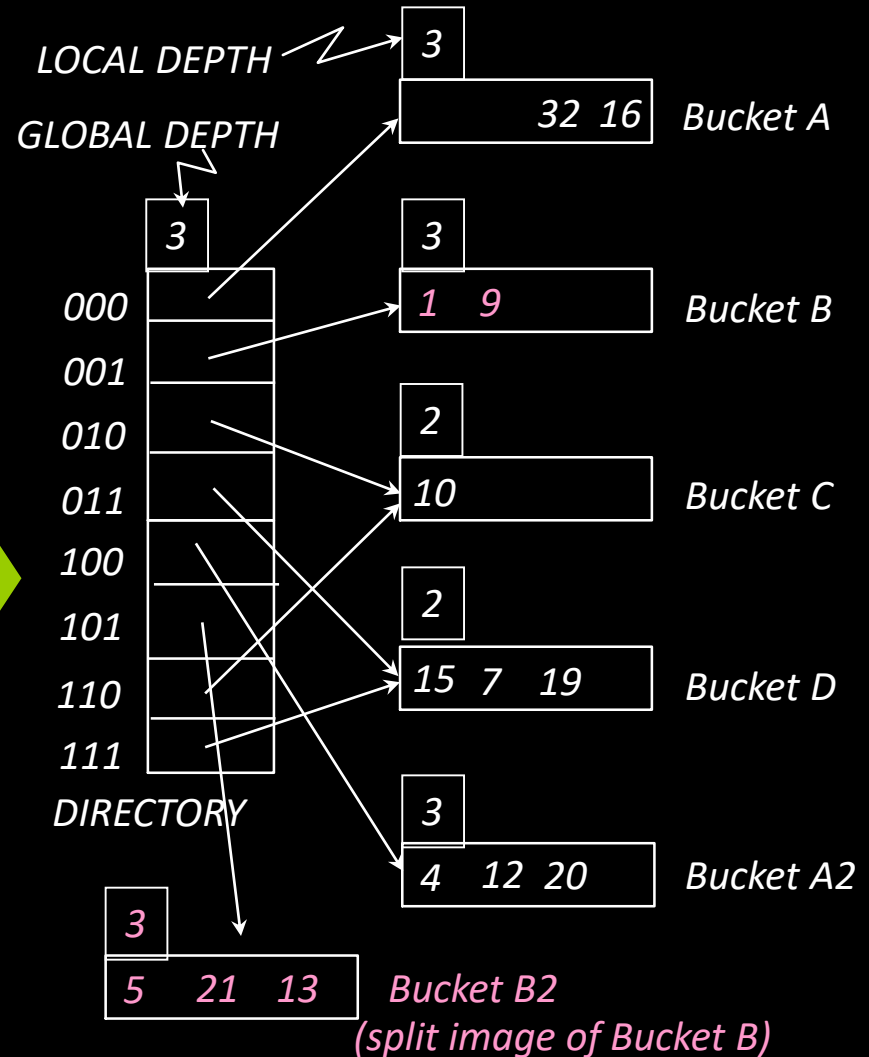

Insert 20 (10100): Double Directory



Split the bucket when local depth = global depth

- **if yes**, double the directory + rehash the entries and distribute into two buckets
- **if no**, rehash the entries and distribute them into two buckets.
- increment the local depth.

Insert 9 (1001): only Split Bucket



Only split bucket:

- Rehash bucket B
- Increment local depth

Extendible Hashing - Points to Note

- What is the global depth of directory?
 - Max # of bits needed to tell which bucket an entry belongs to.
- What is the local depth of a bucket?
 - # of bits used to determine if an entry belongs to a specific bucket
- When does bucket split cause directory doubling?
 - Bucket is full & local depth = global depth.
- How to efficiently double the directory?
 - Directory is doubled by copying it over and 'fixing' pointers to the splitted bucket

Linear Hashing



Linear Hashing (LH)

- Another dynamic hashing scheme, as an alternative to Extendible Hashing.
 - What are problems in static/extendible hashing?
 - Static hashing can cause **long** overflow chains
 - Extendible hashing: data skew causing large **directory**
 - **Data skew** = Multiple entries with same hash value
 - If bucket already full of same hash value, will keep doubling forever!
 - Is it possible to come up with a more balanced solution?
 - Shorter overflow chains than static hashing
 - **No need for directory**
- => Linear Hashing is the answer**

Linear Hashing Algorithm

- Initial Stage.
 - The initial stage distributes entries into N_0 buckets.
 - The hash function to perform is noted h_0
- *Idea*: Use a family of hash functions h_0, h_1, h_2, \dots
 - $h_i(\text{key}) = \text{key} \bmod (2^i N_0)$; N_0 = initial # buckets
 - h_{i+1} doubles the range of h_i (similar to directory doubling)

Example



Linear Hashing Verbal Algorithm


- Splitting buckets
 - If a **bucket overflows** then its primary bucket is chained to an overflow bucket + the bucket **having the next pointer** will get split
 - The bucket to split is the one having the **next pointer** (*not* necessarily the bucket that overflows) – splitting done in round-robin fashion
 - The **next** pointer is moved to the next bucket... and so on until the N^{th} bucket has been split
 - When a bucket is split its entries (including those in overflow buckets) are distributed using h_1
 - To access split buckets the next level hash function (h_1) is applied
 - H_1 maps entries to $2N_0$ buckets
- Alternatively, splitting can be triggered when a fill factor (**average occupancy in buckets** e.g. 80%) is exceeded

Linear Hashing (Contd.)

- Directory avoided in LH by using overflow buckets, and choosing a bucket to split in round-robin.
 - **Splitting proceeds in 'rounds'**. Round ends when all N_R (for round R) initial buckets are split.
 - **Search:** To find bucket for data entry r use $h_{round}(r)$
 - If $h_{round}(r)$ in range '*Next to N_R* ', r belongs here.
 - Else (if it is hashed to bucket **before** Next pointer) then must apply $h_{round+1}(r)$ to find the bucket

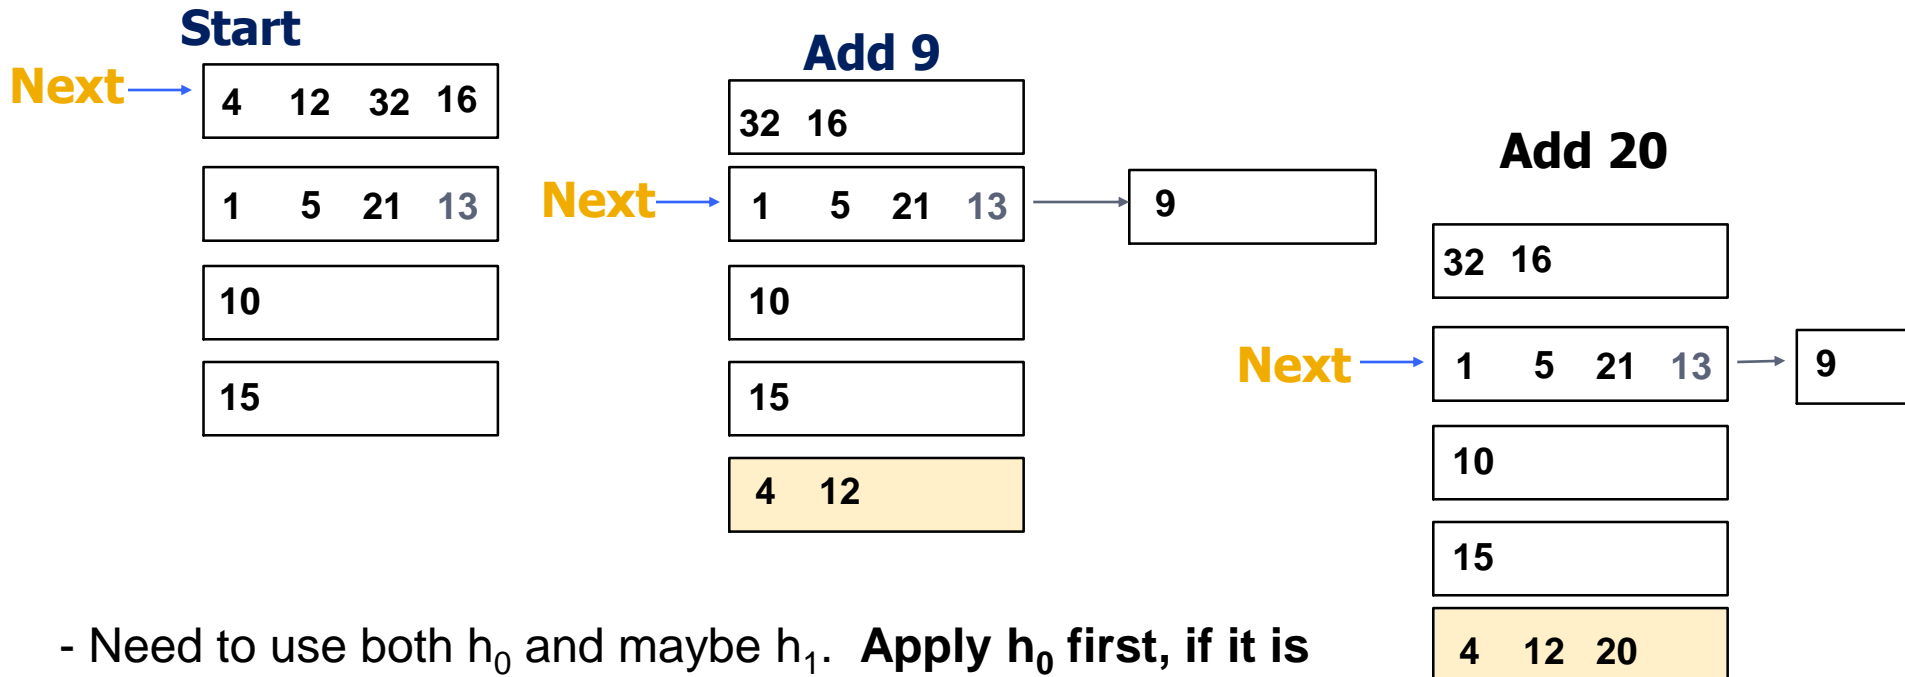
Finding bucket for a search key value

- In the middle of a round.



Linear Hashing Example

- Let's start with $N_0 = 4$ buckets, $H_0 = \text{key mod } 4$
 - Start at round₀ with next pointer pointing to bucket B_0
 - Each time any bucket fills, split "next" bucket and redistribute the values using H_1 ($H_1 = \text{key mod } 8$)




- Need to use both h_0 and maybe h_1 . **Apply h_0 first, if it is hashed to bucket before next pointer then use h_1 .**

If $(h_0(\text{key}) < \text{Next})$ then Use $h_1(\text{key})$ instead


Linear Hashing Example (cont)

- Overflow chains do exist, but eventually get split
- Instead of doubling, new buckets added one-at-a-time


Add 6



Add 17



Linear Hashing Example



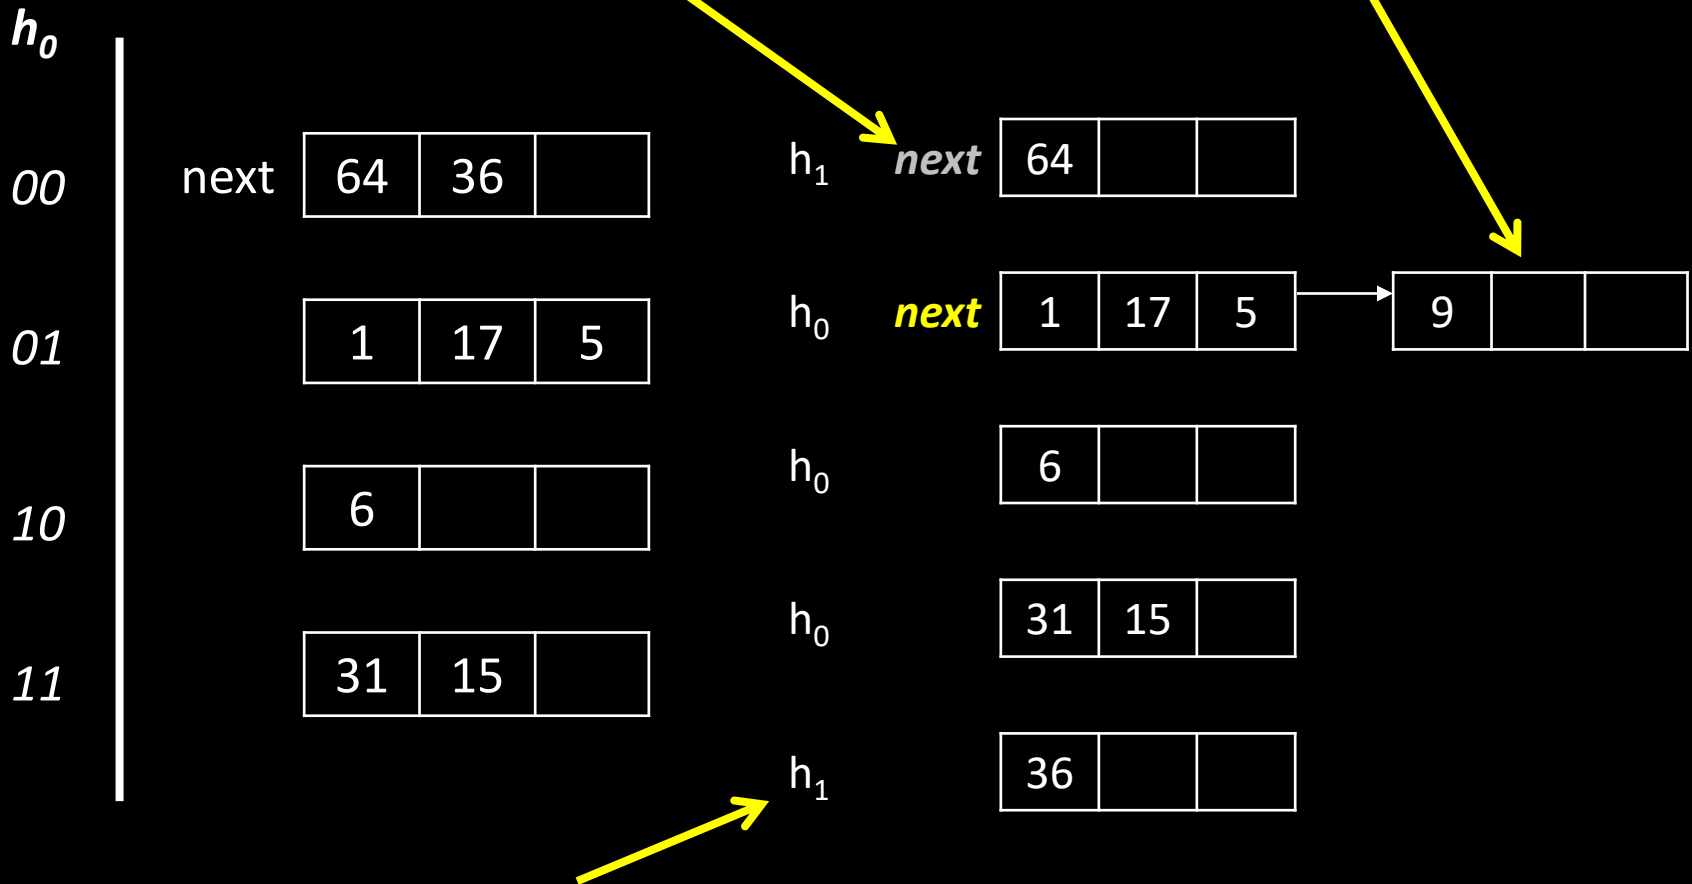
- Initially, we have $N_0 = 4$ buckets
- Assume three entries fit on a bucket
- $h_0 = \text{key mod } 4$
- Next** pointer indicates which bucket is to split next (Round Robin)
- Now consider what happens when 9 (1001) is inserted (which will not fit in the second bucket)

Insert 9 (1001)

- The bucket indicated by **next** (the first one) is split
- Next** is incremented.


An overflow bucket is chained to the primary bucket to contain the inserted value.

This causes a split to occur.




For subsequent inserts, apply hash function h_0 first, if the key is hashed to bucket before **next** pointer. Then h_1 must be used to insert the new entry.

Insert 8 (1000), 7(111), 18(10010), 14(1100)



- Note that the split bucket is not necessarily the overflow bucket. The split bucket is chosen based on round robin.
- Need to use both h_0 and maybe h_1 . **Apply h_0 first, if it is hashed to bucket before next pointer. Then use h_1 .**

Before Insert 11 (1011)



After Insert 11 (1011)

Which hash function (h0 or h1) for searching 9?

How about searching for 18?

(000) H_1	B_0	64	8	
-------------	-------	----	---	--

(001) H_1	<i>next</i> B_1	1	17	9
-------------	-------------------	---	----	---

(10) H_0	<i>next</i> B_2	6	18	14
------------	-------------------	---	----	----

(11) H_0	B_3				\longrightarrow			
		31	15	7		11		

(100) H_1	B_4	36		
-------------	-------	----	--	--

(101) H_1	B_5	5		
-------------	-------	---	--	--

Inserting 11 caused an overflow of B_3 .
This triggered splitting B_1

Insert 32, 16

(000) H_1 B_0

64	8	32
----	---	----

 \longrightarrow

16		
----	--	--

(001) H_1 B_1

1	17	9
---	----	---

(010) H_1 *next* B_2

18		
----	--	--

(11) H_0 *next* B_3

31	15	7
----	----	---

 \longrightarrow

11		
----	--	--

(100) H_1 B_4

36		
----	--	--

(101) H_1 B_5

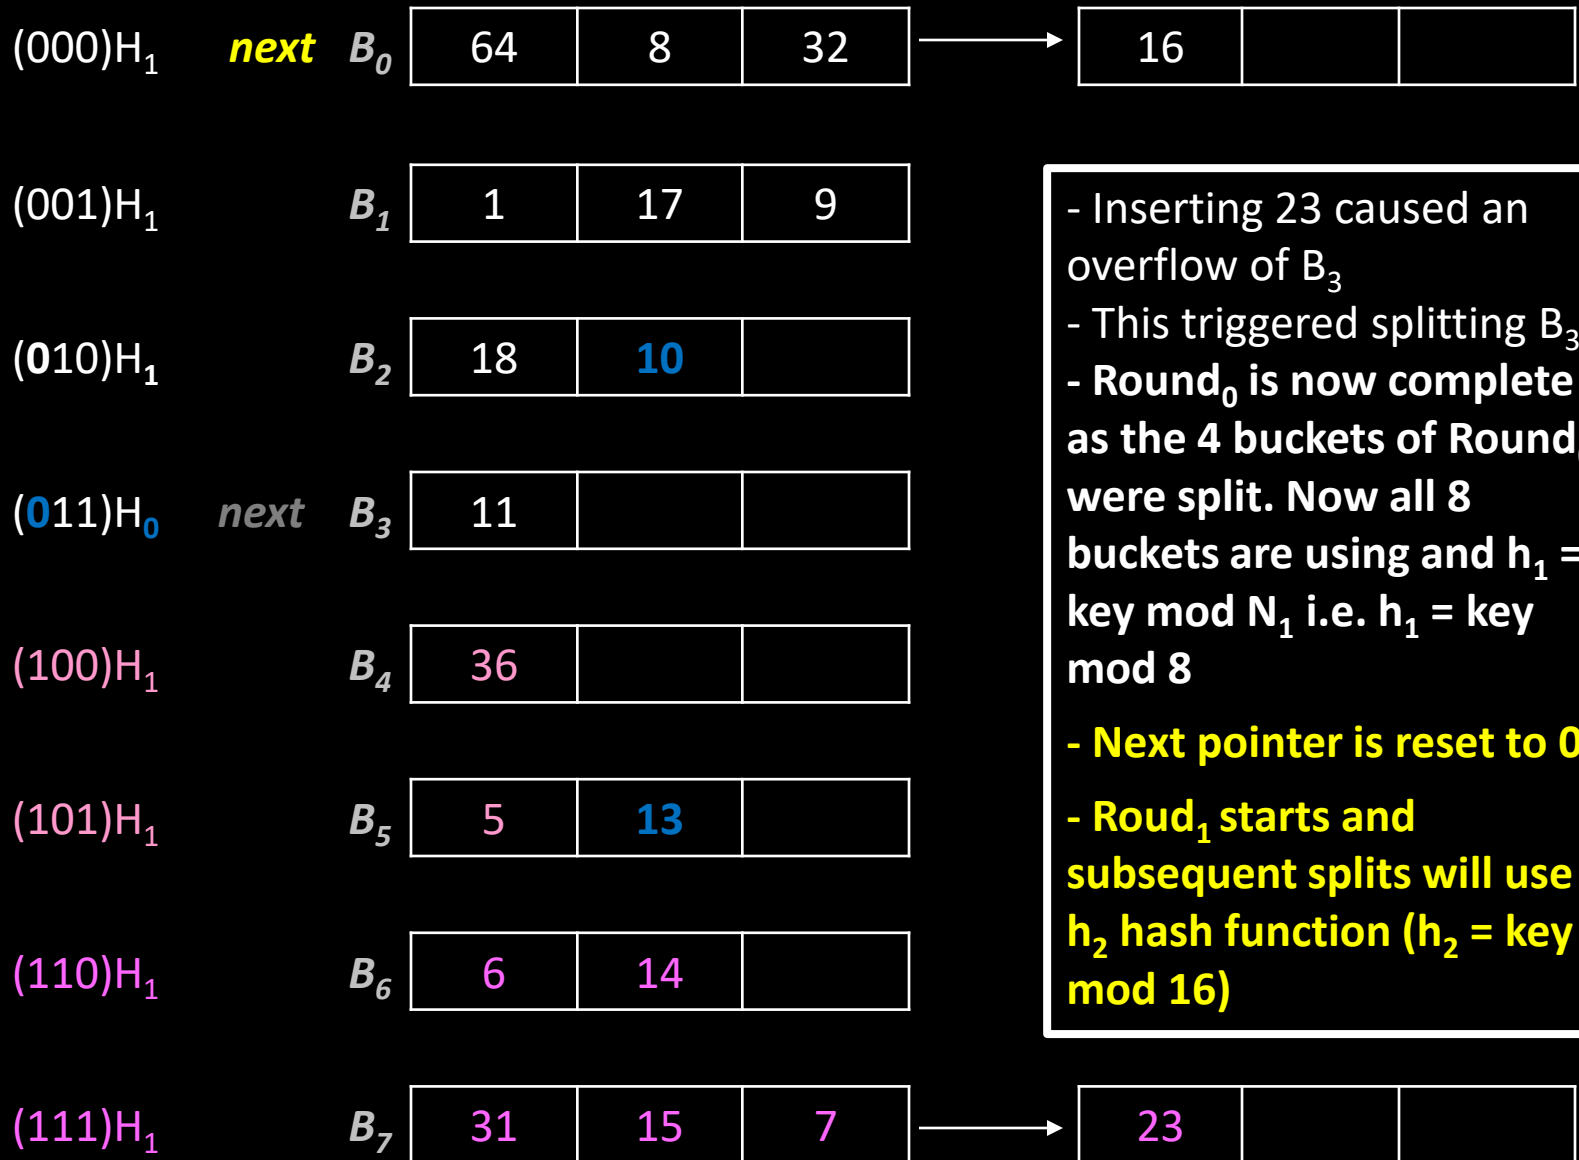
5		
---	--	--

(110) H_1 B_6

6	14	
---	----	--

Inserting 16 caused an overflow of B_0 .
This triggered splitting B_2

Insert 10, 13, 23



- Inserting 23 caused an overflow of B_3
- This triggered splitting B_3
- **Round₀ is now complete as the 4 buckets of Round₀ were split. Now all 8 buckets are using and $h_1 = \text{key mod } N_1$ i.e. $h_1 = \text{key mod } 8$**
- **Next pointer is reset to 0**
- **Round₁ starts and subsequent splits will use h_2 hash function ($h_2 = \text{key mod } 16$)**

Extendible and Linear Hashing


- Linear hashing does not require a dictionary
- Linear hashing may result in less space efficiency because buckets are split before they overflow
- Multiple collisions in one bucket in extensible hashing will result in a large directory
 - Such a directory may not fit on one disk block
- Collisions in linear hashing lead to long overflow chains for the bucket with the collisions
 - Requiring multiple disk reads for that bucket
 - But no increase in the cost of accessing other buckets

Summary

- Use a hash index for point queries only.
Use a B-tree for multipoint queries or range queries
- Use clustered index:
 - if your queries need all or most of the table columns
 - if multipoint, range queries or order-by queries are often requested

Appendix - Cryptographic Hashing

Cryptographic Hash Function Analogy



$h = H(m)$. Hash
variable size
message m to
produce a **fixed size**
hash value (sometimes
called a *message digest*)

Cryptographic Hash Functions

- MD5 (Message Digest 5)
 - Produces a 128-bit hash
 - Collisions can be found. An attacker can use them to substitute an authorized message with an unauthorized one.
- SHA1 (Secure Hash Algorithm 1)
 - 160-bit hash
 - Collisions can be found
- SHA2
 - Actually 4 different hash functions: SHA-224, SHA-256, SHA-384, SHA-512
 - Minor attacks, but still good
- SHA3
 - New NIST standard
 - No known attacks

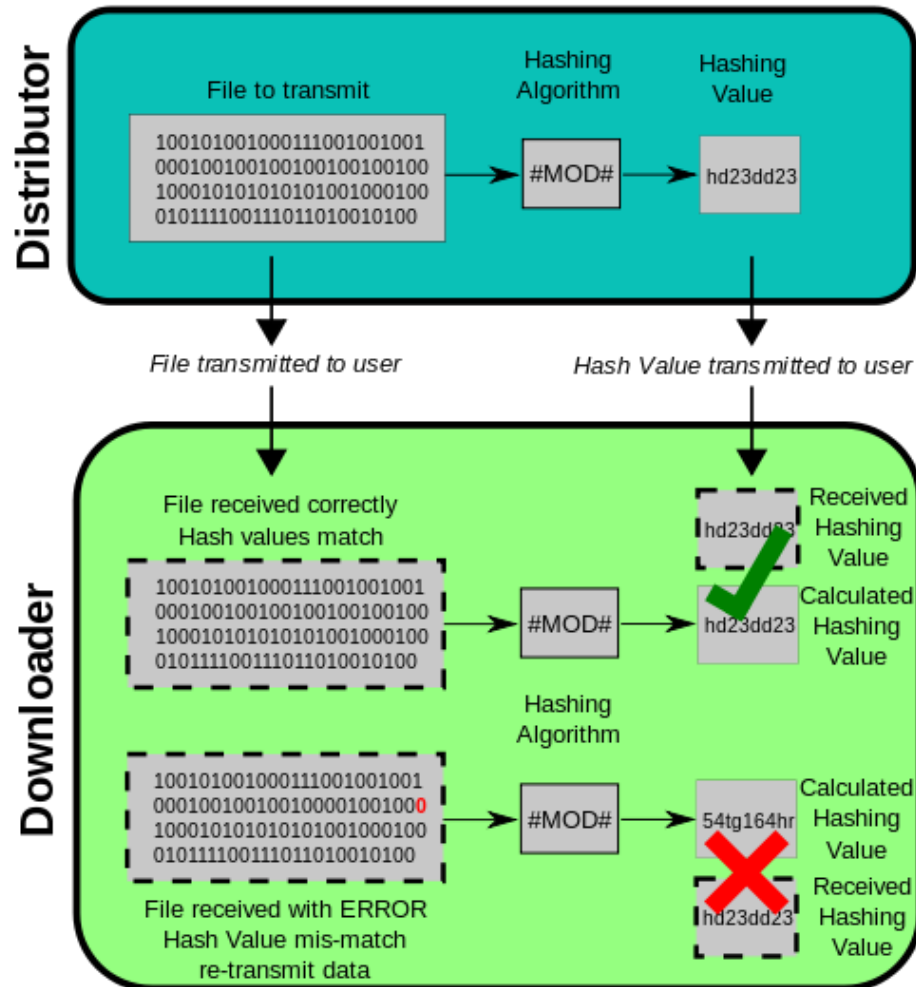
Sample in Python



```
import hashlib
md = hashlib.sha256(b"The quick brown fox jumps over the lazy dog").hexdigest()
print (md)
```

d7a8fbb307d7809469ca9abcb0082e4f8d5651e46d3cdb762d02d0bf37c9e592

Application 1: File Transmission



Application 2: Message Authentication Code (MAC)

