

Hashing

CMPT 6o6



Read Chapter 18 and Chapter
17 (pages 611 to 621)

Outline


- ① Extensible Hashing
- ② Linear Hashing

Hash Indexes



Hash Tables in Main Memory

- In a hash table a hash function maps search key values to array elements
 - The array can either contain the data objects, or
 - Linked lists containing data objects (often called *buckets*)
- Hash functions generate a value between 0 and $B-1$
 - Where B is the number of buckets
 - A record with **search key K** is **stored in bucket $h(K)$**

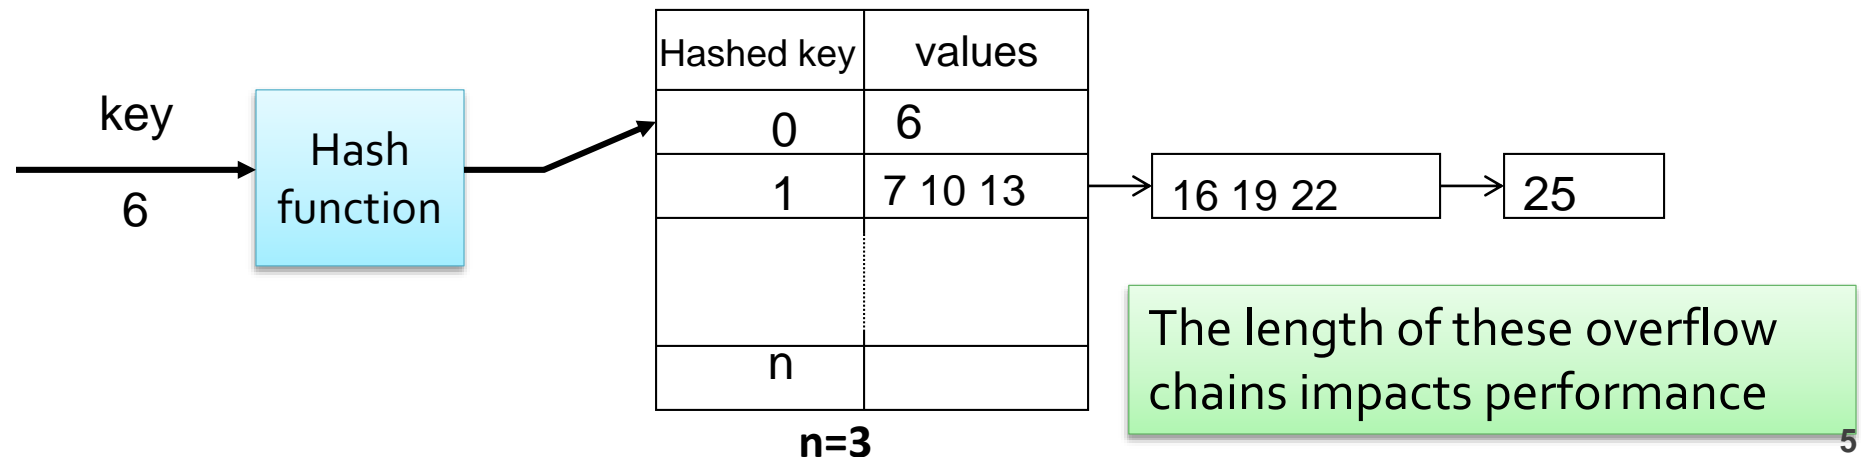


- A typical hash function is a bit representation of (the Search Key Value **modulus** the Number of Buckets)
- Hash indexes do not support range lookup

Hash Index



- A hash index stores **key-value pairs** based on a *hash function* => *an array of pointers to buckets*
- The hash function often is **$K \bmod B$**
 - Where **K** is the key value and **B** is the number of buckets
- Collision is resolved by placing new values in an **overflow bucket**



Static Hashing vs. Dynamic Hashing

- In static hashing the number of buckets is fixed
 - If the file grows most buckets will have overflow chains, reducing efficiency

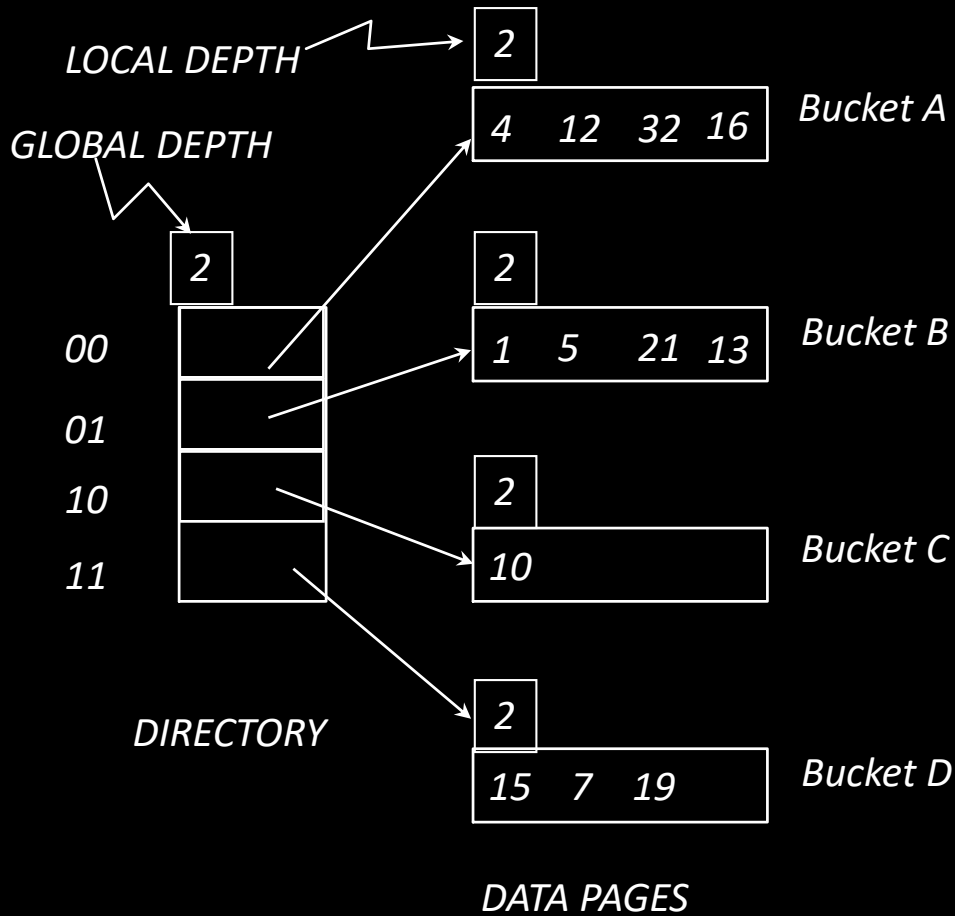
=> There should be enough buckets so that there are few overflow blocks
- Two versions of dynamic hashing that allow **dynamic File Expansion**
 - Extensible hashing
 - Linear hashing

Extensible Hashing Structure

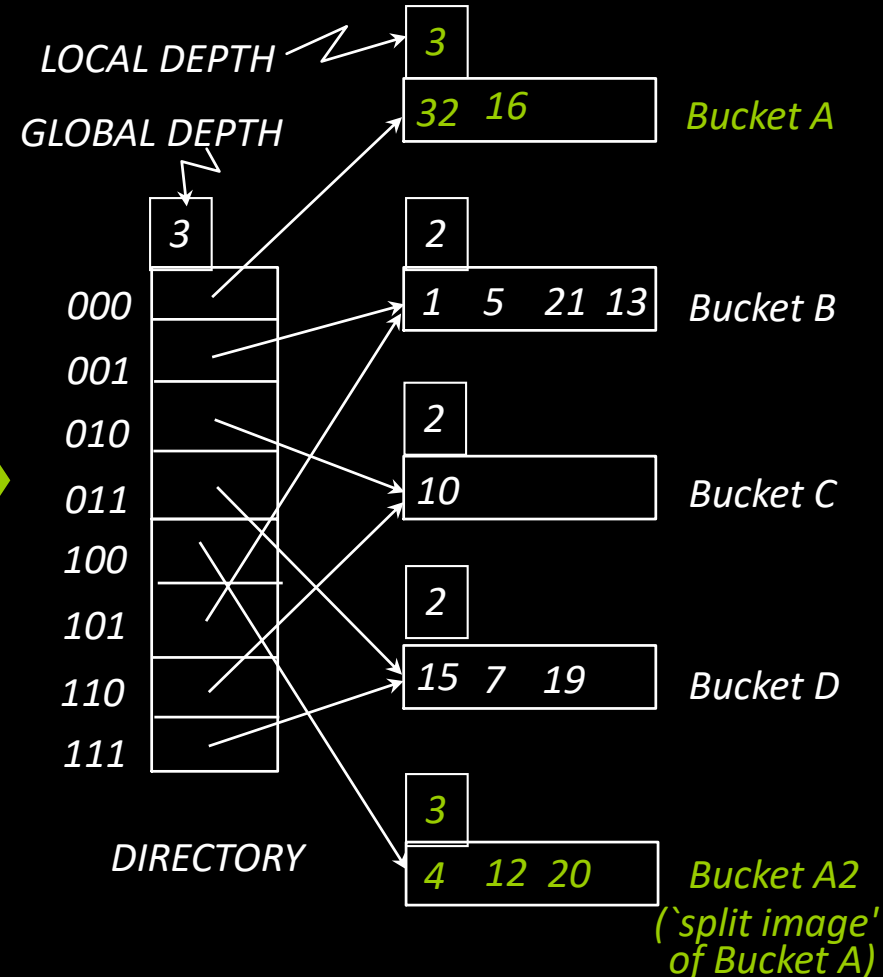
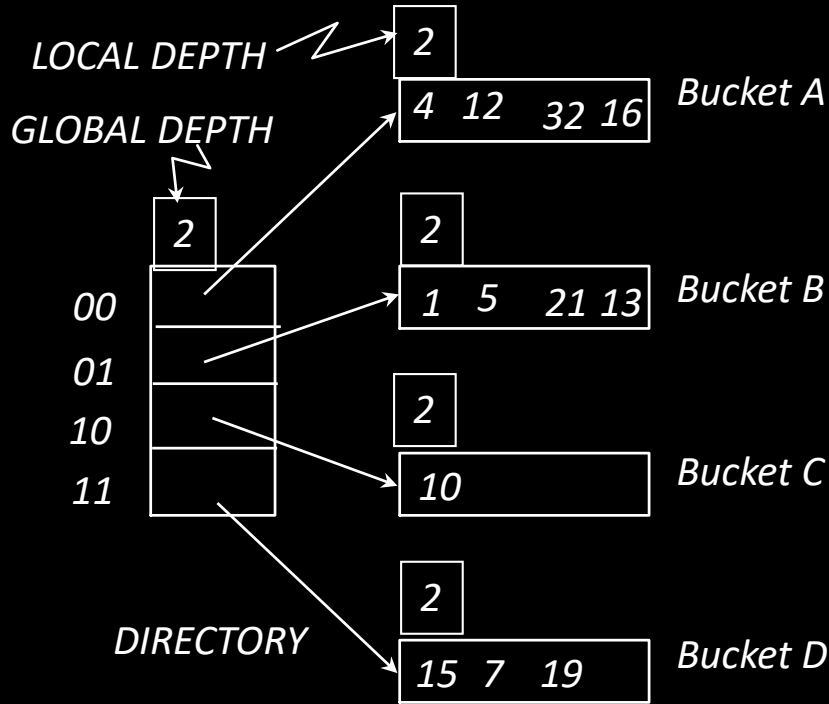
- In extensible hashing there is a **directory** to buckets
 - The directory is an array of pointers to buckets
 - As the array only contains pointers it is relatively small, so it can usually fit in memory
 - The directory size is doubled when overflow occurs
 - New buckets are only created as necessary
- The hash function computes a sequence of bits
 - The directory (and the associated buckets) **uses the last i bits** of the hash value
 - The directory will have **2^i** entries
 - When the directory grows, **$i+1$** bits of the hash value are used

Extensible Hashing Example

- Directory is array of size 4.
- Assume a bucket capacity of 4
- What info this hash data structure maintains?
 - **Global Depth**: # of bits needed to tell which bucket an entry belongs to
 - **Local Depth**: # of bits used to determine if an entry belongs to a specific bucket



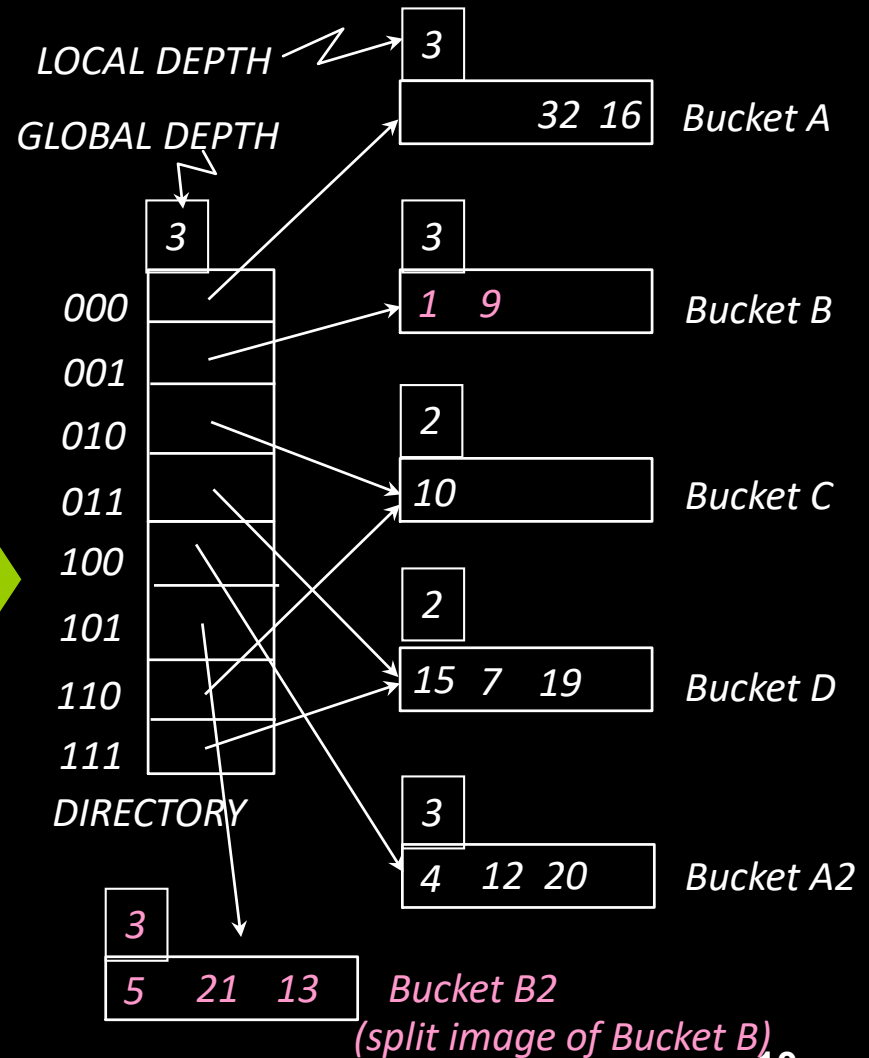

Insert 20 (10100): Double Directory



Split the bucket when local depth = global depth

- **if yes**, double the directory + rehash the entries and distribute into two buckets
- **if no**, rehash the entries and distribute them into two buckets.
- increment the local depth.

Insert 9 (1001): only Split Bucket



Only split bucket:

- Rehash bucket B
- Increment local depth

Extensible Hashing - Points to Note

- What is the global depth of directory?
 - Max # of bits needed to tell which bucket an entry belongs to.
- What is the local depth of a bucket?
 - # of bits used to determine if an entry belongs to a specific bucket
- When does bucket split cause directory doubling?
 - Bucket is full & local depth = global depth.
- How to efficiently double the directory?
 - Directory is doubled by copying it over and 'fixing' pointers to the splitted bucket



Linear Hashing (LH)

- Another dynamic hashing scheme, as an alternative to Extendible Hashing.
 - What are problems in static/extendible hashing?
 - Static hashing can cause **long** overflow chains
 - Extendible hashing: data skew causing large **directory**
 - **Data skew** = Multiple entries with same hash value
 - If bucket already full of same hash value, will keep doubling forever!
 - Is it possible to come up with a more balanced solution?
 - Shorter overflow chains than static hashing
 - **No need for directory**
- => Linear Hashing is the answer**

Linear Hashing Algorithm

- Initial Stage.
 - The initial level distributes entries into N_0 buckets.
 - The hash function to perform is noted h_0
- *Idea*: Use a family of hash functions h_0, h_1, h_2, \dots
 - $h_i(\text{key}) = \text{key} \bmod(2^i N_0)$; N_0 = initial # buckets
 - h_{i+1} doubles the range of h_i (similar to directory doubling)

Example




Linear Hashing Verbal Algorithm

- Splitting buckets
 - If a **bucket overflows** its primary page is chained to an overflow page + **some** bucket is split
 - The first bucket to be split is the first bucket (*not* necessarily the bucket that overflows)
 - The **next** bucket to be split is the second bucket in the file ... and so on until the N^{th} has been split
 - When buckets are split their entries (including those in overflow pages) are distributed using h_1
 - To access split buckets the next level hash function (h_1) is applied
 - H_1 maps entries to $2N_0$ buckets
- Alternatively, splitting can be triggered when a fill factor (**average occupancy in buckets** e.g. 80%) is exceeded


Linear Hashing Example

- Let's start with $N = 4$ Buckets
 - Start at "round" 0, "next" 0
 - Each time any bucket fills, split "next" bucket


Start



Add 9



Add 20




If $(0 \leq h_{\text{next}}(\text{key}) < N_{\text{next}})$ then Use $h_{\text{next}}(\text{key})$ instead


Linear Hashing Example (cont)

- Overflow chains do exist, but eventually get split
- Instead of doubling, new buckets added one-at-a-time

Add 6



Add 17




Linear Hashing (Contd.)

- Directory avoided in LH by using overflow buckets, and choosing bucket to split in round-robin.
 - **Splitting proceeds in 'rounds'**. Round ends when all N_R initial (for round R) buckets are split.
 - **Search:** To find bucket for data entry r , find $h_{round}(r)$:
 - If $h_{round}(r)$ in range '*Next to N_R* ', r belongs here.
 - Else, must apply $h_{round+1}(r)$ to find out.

Overview of LH File

- In the middle of a round.



Linear Hashing Example

h_0				
00	next <table><tr><td>64</td><td>36</td><td></td></tr></table>	64	36	
64	36			
01	<table><tr><td>1</td><td>17</td><td>5</td></tr></table>	1	17	5
1	17	5		
10	<table><tr><td>6</td><td></td><td></td></tr></table>	6		
6				
11	<table><tr><td>31</td><td>15</td><td></td></tr></table>	31	15	
31	15			

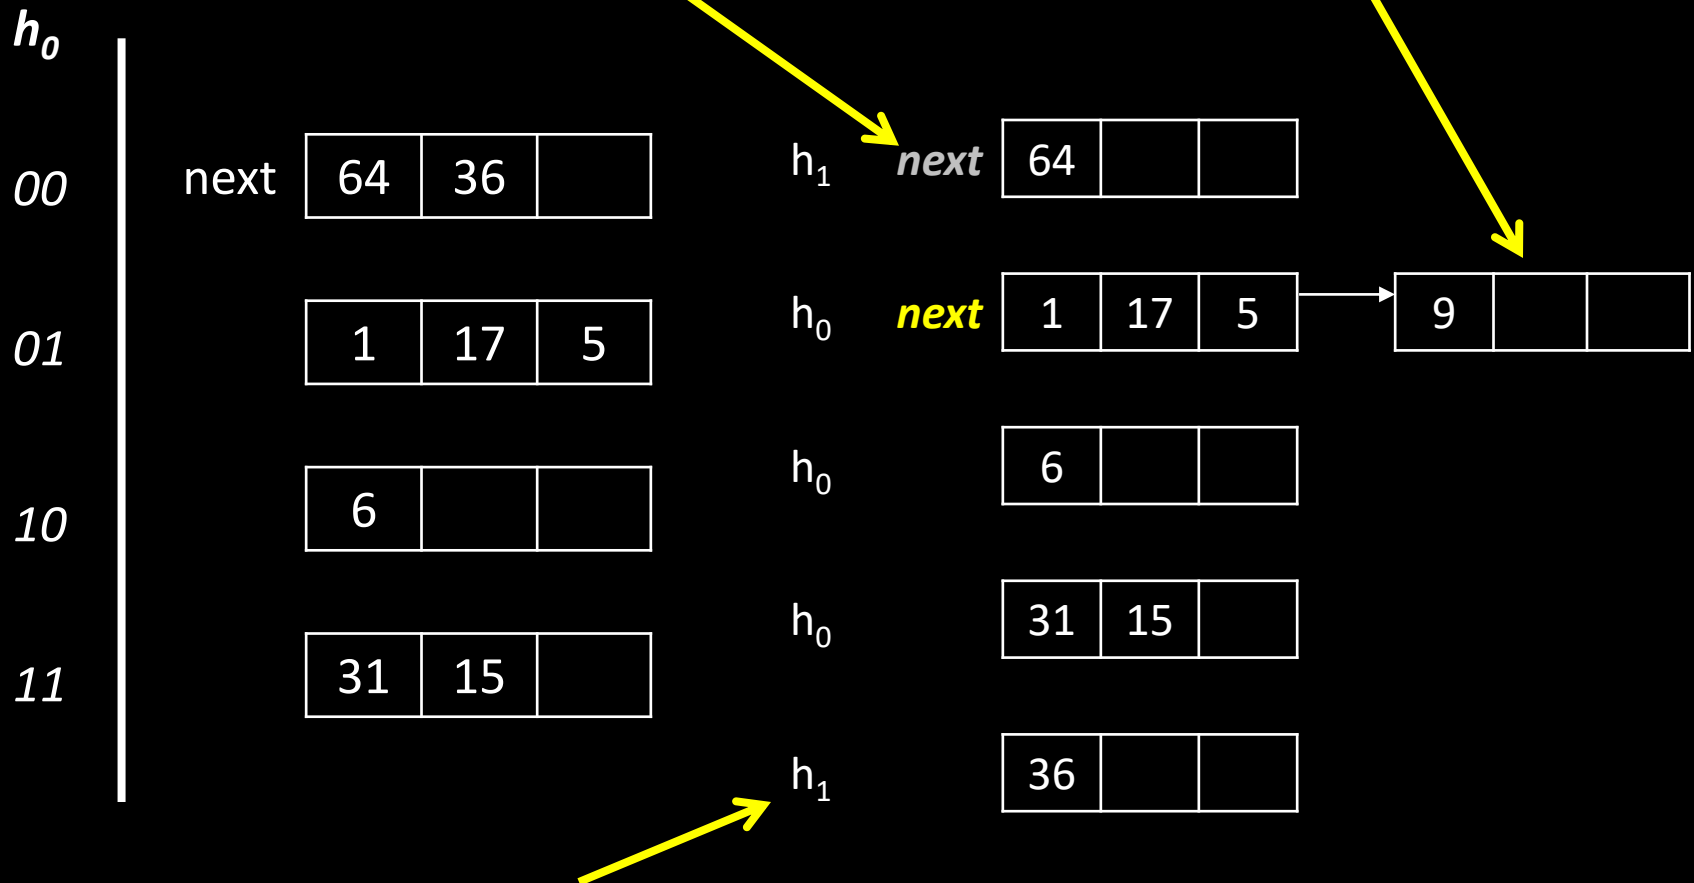
- Initially, we have $N_0 = 4$ buckets
- Assume three entries fit on a bucket
- h_0 range = 4 buckets
- Note that **next** indicates which bucket is to split next (Round Robin)
- Now consider what happens when 9 (1001) is inserted (which will not fit in the second bucket)

Insert 9 (1001)

- The bucket indicated by **next** (the first one) is split
- **Next** is incremented.


An overflow bucket is chained to the primary bucket to contain the inserted value.

This causes a split to occur.




For subsequent inserts, apply hash function h_0 first, if the key is hashed to bucket before **next** pointer. Then h_1 must be used to insert the new entry.

Insert 8 (1000), 7(111), 18(10010), 14(1100)



- Note that the split bucket is not necessary the overflow bucket. The split bucket is chosen based on round robin.
- Need to use both h_0 and maybe h_1 . **Apply h_0 first, if it is hashed to bucket before next pointer. Then use h_1 .**

Before Insert 11 (1011)



After Insert 11 (1011)

Which hash function
(h0 or h1) for
searching 9?

(000)H ₁	64	8	
---------------------	----	---	--

(001)H ₁	<i>next</i>	1	17	9
---------------------	-------------	---	----	---

(10)H ₀	<i>next</i>	6	18	
--------------------	-------------	---	----	--

(11)H ₀	31	15	7	→	11		
--------------------	----	----	---	---	----	--	--

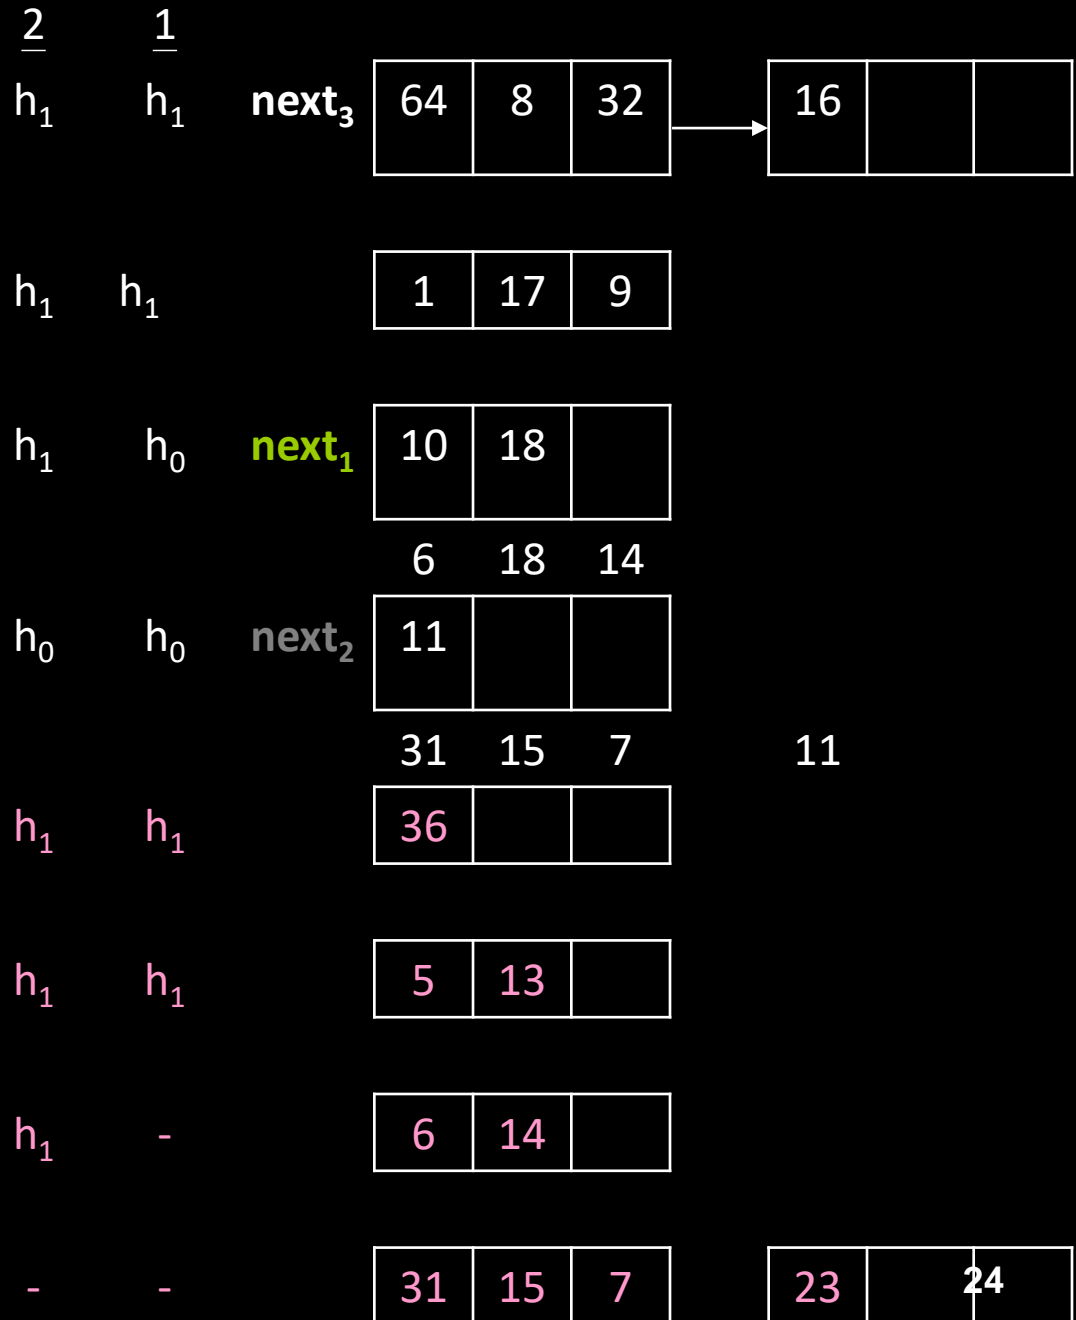
(100)H ₁	36	14	
---------------------	----	----	--

(101)H ₁	5		
---------------------	---	--	--

How about searching
for 18?

Linear Hashing

- Insert 32, 16₂
- Insert 10, 13, 23₃
- After the 2nd split the base level is 1 ($N_1 = 8$), use h_1 .
- Subsequent splits will use h_2 for inserts between the first bucket and *next*-1.



Extensible and Linear Hashing

- Linear hashing does not require a dictionary
- Linear hashing may result in less space efficiency because buckets are split before they overflow
- Multiple collisions in one bucket in extensible hashing will result in a large directory
 - Such a directory may not fit on one disk block
- Collisions in linear hashing lead to long overflow chains for the bucket with the collisions
 - Requiring multiple disk reads for that bucket
 - But no increase in the cost of accessing other buckets

Summary

- Use a hash index for point queries only.
Use a B-tree if multipoint queries or range queries are used
- Use clustering:
 - if your queries need all or most of the fields of each records returned
 - if multipoint, range queries or order-by queries are often requested